

Area Plots, Histograms, and Bar Charts

Estimated time needed: **30** minutes

Objectives

After completing this lab you will be able to:

- Create and customize
 - Area plot
 - Histogram
 - Bar charts on a dataset

Table of Contents

Import Libraries

Import the `matplotlib` library.

```
#Import Primary Modules:
import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library

# use the inline backend to generate the plots within the browser
%matplotlib inline

import matplotlib as mpl
import matplotlib.pyplot as plt

mpl.style.use('ggplot') # optional: for ggplot-like style

# check for latest version of Matplotlib
print('Matplotlib version: ', mpl.__version__) # >= 2.0.0

Matplotlib version: 3.5.3
```

Fetching Data

Dataset: Immigration to Canada from 1980 to 2013 - [International migration flows to and from selected countries - The 2015 revision](#) from United Nation's website

In this lab, we will focus on the Canadian Immigration data and use the **already cleaned dataset** and can be fetched from here. You can refer to the lab on data pre-processing wherein this dataset is cleaned for a quick refresh your Panads skills [Data pre-processing with Pandas](#)

```
df_can = pd.read_csv('https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DV0101EN-SkillsNetwork/Data%20Files/Canada.csv')
```

```
print('Data read into a pandas dataframe!')
```

Data read into a pandas dataframe!

Let's take a look at the first five items in our dataset.

```
df_can.head()
```

	Country	Continent	Region	DevName	1980
1981 \					
0	Afghanistan	Asia	Southern Asia	Developing regions	16
39					
1	Albania	Europe	Southern Europe	Developed regions	1
0					
2	Algeria	Africa	Northern Africa	Developing regions	80
67					
3	American Samoa	Oceania	Polynesia	Developing regions	0
1					
4	Andorra	Europe	Southern Europe	Developed regions	0
0					

	1982	1983	1984	1985	...	2005	2006	2007	2008	2009	2010
2011 \											
0	39	47	71	340	...	3436	3009	2652	2111	1746	1758
2203											
1	0	0	0	0	...	1223	856	702	560	716	561
539											
2	71	69	63	44	...	3626	4807	3623	4005	5393	4752
4325											
3	0	0	0	0	...	0	1	0	0	0	0
0											
4	0	0	0	0	...	0	1	1	0	0	0
0											

	2012	2013	Total
0	2635	2004	58639
1	620	603	15699
2	3774	4331	69439
3	0	0	6
4	1	1	15

[5 rows x 39 columns]

Let's find out how many entries there are in our dataset.

```
# print the dimensions of the dataframe
print(df_can.shape)

(195, 39)
```

Set the country name as index - useful for quickly looking up countries using .loc method.

```
df_can.set_index('Country', inplace=True)

# Let's view the first five elements and see how the dataframe was
changed
df_can.head()
```

	Continent		Region				DevName		1980
1981 \ Country									
Afghanistan 39	Asia		Southern Asia				Developing regions		16
Albania 0	Europe		Southern Europe				Developed regions		1
Algeria 67	Africa		Northern Africa				Developing regions		80
American Samoa 1	Oceania		Polynesia				Developing regions		0
Andorra 0	Europe		Southern Europe				Developed regions		0
	1982	1983	1984	1985	1986	...	2005	2006	2007
2008 \ Country						...			
Afghanistan 2111	39	47	71	340	496	...	3436	3009	2652
Albania 560	0	0	0	0	1	...	1223	856	702
Algeria 4005	71	69	63	44	69	...	3626	4807	3623
American Samoa 0	0	0	0	0	0	...	0	1	0
Andorra 0	0	0	0	0	2	...	0	1	1
	2009	2010	2011	2012	2013	Total			
Country									
Afghanistan	1746	1758	2203	2635	2004	58639			
Albania	716	561	539	620	603	15699			
Algeria	5393	4752	4325	3774	4331	69439			

American Samoa	0	0	0	0	0	6
Andorra	0	0	0	1	1	15

[5 rows x 38 columns]

Notice now the country names now serve as indices.

```
print('data dimensions:', df_can.shape)
```

```
data dimensions: (195, 38)
```

```
# finally, let's create a list of years from 1980 - 2013  
# this will come in handy when we start plotting the data
```

```
years = list(map(str, range(1980, 2014)))
```

```
years
```

```
['1980',  
'1981',  
'1982',  
'1983',  
'1984',  
'1985',  
'1986',  
'1987',  
'1988',  
'1989',  
'1990',  
'1991',  
'1992',  
'1993',  
'1994',  
'1995',  
'1996',  
'1997',  
'1998',  
'1999',  
'2000',  
'2001',  
'2002',  
'2003',  
'2004',  
'2005',  
'2006',  
'2007',  
'2008',  
'2009',  
'2010',  
'2011',
```

```
'2012',  
'2013']
```

Area Plots

In the last module, we created a line plot that visualized the top 5 countries that contributed the most immigrants to Canada from 1980 to 2013. With a little modification to the code, we can visualize this plot as a cumulative plot, also known as a **Stacked Line Plot** or **Area plot**.

```
df_can.sort_values(['Total'], ascending=False, axis=0, inplace=True)
```

```
# get the top 5 entries
```

```
df_top5 = df_can.head()
```

```
# transpose the dataframe
```

```
df_top5 = df_top5[years].transpose()
```

```
df_top5.head()
```

Country	India	China	United Kingdom of Great Britain and Northern Ireland
1980	8880	5123	
22045			
1981	8670	6682	
24796			
1982	8147	3308	
20620			
1983	7338	1863	
10015			
1984	5704	1527	
10170			

Country	Philippines	Pakistan
1980	6051	978
1981	5921	972
1982	5249	1201
1983	4562	900
1984	3801	668

Area plots are stacked by default. And to produce a stacked area plot, each column must be either all positive or all negative values (any `NaN`, i.e. not a number, values will default to 0). To produce an unstacked plot, set parameter `stacked` to value `False`.

```
# let's change the index values of df_top5 to type integer for plotting
```

```
df_top5.index = df_top5.index.map(int)
```

```
df_top5.plot(kind='area',
```


```

        stacked=False,
        figsize=(20, 10)) # pass a tuple (x, y) size

plt.title('Immigration Trend of Top 5 Countries')
plt.ylabel('Number of Immigrants')
plt.xlabel('Years')

plt.show()

```



The unstacked plot has a default transparency (alpha value) at 0.5. We can modify this value by passing in the `alpha` parameter.


```

df_top5.plot(kind='area',
              alpha=0.25, # 0 - 1, default value alpha = 0.5
              stacked=False,
              figsize=(20, 10))

plt.title('Immigration Trend of Top 5 Countries')
plt.ylabel('Number of Immigrants')
plt.xlabel('Years')

plt.show()

```



Two types of plotting

As we discussed in the video lectures, there are two styles/options of plotting with `matplotlib`, plotting using the Artist layer and plotting using the scripting layer.

Option 1: Scripting layer (procedural method) - using `matplotlib.pyplot` as 'plt'

You can use `plt` i.e. `matplotlib.pyplot` and add more elements by calling different methods procedurally; for example, `plt.title(...)` to add title or `plt.xlabel(...)` to add label to the x-axis.

```
# Option 1: This is what we have been using so far
df_top5.plot(kind='area', alpha=0.35, figsize=(20, 10))
plt.title('Immigration trend of top 5 countries')
plt.ylabel('Number of immigrants')
plt.xlabel('Years')
```

Option 2: Artist layer (Object oriented method) - using an `AXES` instance from `Matplotlib` (preferred)

You can use an `AXES` instance of your current plot and store it in a variable (eg. `ax`). You can add more elements by calling methods with a little change in syntax (by adding "`set_`" to the previous methods). For example, use `ax.set_title()` instead of `plt.title()` to add title, or `ax.set_xlabel()` instead of `plt.xlabel()` to add label to the x-axis.


This option sometimes is more transparent and flexible to use for advanced plots (in particular when having multiple plots, as you will see later).

In this course, we will stick to the **scripting layer**, except for some advanced visualizations where we will need to use the **artist layer** to manipulate advanced aspects of the plots.

```
# option 2: preferred option with more flexibility
ax = df_top5.plot(kind='area', alpha=0.35, figsize=(20, 10))

ax.set_title('Immigration Trend of Top 5 Countries')
ax.set_ylabel('Number of Immigrants')
ax.set_xlabel('Years')

Text(0.5, 0, 'Years')
```



Question: Use the scripting layer to create a stacked area plot of the 5 countries that contributed the least to immigration to Canada **from** 1980 to 2013. Use a transparency value of 0.45.

```
### type your answer here
df_least5 = df_can.tail(5)
df_least5 = df_least5[years].transpose()
df_least5.head()
```

Country	San Marino	New Caledonia	Marshall Islands	Western Sahara
Palau				
1980	1	0	0	0
0				
1981	0	0	0	0
0				
1982	0	0	0	0
0				
1983	0	0	0	0
0				
1984	0	0	0	0
0				


```

# get the 5 countries with the least contribution
df_least5 = df_can.tail(5)


# transpose the dataframe
df_least5 = df_least5[years].transpose()
df_least5.head()

df_least5.index = df_least5.index.map(int) # let's change the index
values of df_least5 to type integer for plotting
df_least5.plot(kind='area', alpha=0.45, figsize=(20, 10))

plt.title('Immigration Trend of 5 Countries with Least Contribution to
Immigration')
plt.ylabel('Number of Immigrants')
plt.xlabel('Years')

plt.show()

```



Question: Use the artist layer to create an unstacked area plot of the 5 countries that contributed the least to immigration to Canada **from** 1980 to 2013. Use a transparency value of 0.55.

```

# get the 5 countries with the least contribution
df_least5 = df_can.tail(5)

# transpose the dataframe
df_least5 = df_least5[years].transpose()

df_least5.head()

```

```


df_least5.index = df_least5.index.map(int) # let's change the index
values of df_least5 to type integer for plotting

ax = df_least5.plot(kind='area', alpha=0.55, stacked=False,
figsize=(20, 10))

ax.set_title('Immigration Trend of 5 Countries with Least Contribution
to Immigration')
ax.set_ylabel('Number of Immigrants')
ax.set_xlabel('Years')

Text(0.5, 0, 'Years')

```



Histograms

A histogram is a way of representing the *frequency* distribution of numeric dataset. The way it works is it partitions the x-axis into *bins*, assigns each data point in our dataset to a bin, and then counts the number of data points that have been assigned to each bin. So the y-axis is the frequency or the number of data points in each bin. Note that we can change the bin size and usually one needs to tweak it so that the distribution is displayed nicely.

Question: What is the frequency distribution of the number (population) of new immigrants from the various countries to Canada in 2013?

Before we proceed with creating the histogram plot, let's first examine the data split into intervals. To do this, we will use **Numpy's** `histogram` method to get the bin ranges and frequency counts as follows:

```
# let's quickly view the 2013 data
df_can['2013'].head()

Country
India                                     33087
China                                    34129
United Kingdom of Great Britain and Northern Ireland    5827
Philippines                                   29544
Pakistan                                   12603
Name: 2013, dtype: int64

# np.histogram returns 2 values
count, bin_edges = np.histogram(df_can['2013'])

print(count) # frequency count
print(bin_edges) # bin ranges, default = 10 bins

[178  11   1   2   0   0   0   0   1   2]
[  0.  3412.9  6825.8 10238.7 13651.6 17064.5 20477.4 23890.3
27303.2
30716.1 34129. ]
```

By default, the `histogram` method breaks up the dataset into 10 bins. The figure below summarizes the bin ranges and the frequency distribution of immigration in 2013. We can see that in 2013:


- 178 countries contributed between 0 to 3412.9 immigrants
- 11 countries contributed between 3412.9 to 6825.8 immigrants
- 1 country contributed between 6825.8 to 10238.7 immigrants, and so on..

We can easily graph this distribution by passing `kind=hist` to `plot()`.

```
df_can['2013'].plot(kind='hist', figsize=(8, 5))

# add a title to the histogram
plt.title('Histogram of Immigration from 195 Countries in 2013')
# add y-label
plt.ylabel('Number of Countries')
# add x-label
plt.xlabel('Number of Immigrants')

plt.show()
```



In the above plot, the x-axis represents the population range of immigrants in intervals of 3412.9. The y-axis represents the number of countries that contributed to the aforementioned population.

Notice that the x-axis labels do not match with the bin size. This can be fixed by passing in a `xticks` keyword that contains the list of the bin sizes, as follows:

```
# 'bin_edges' is a list of bin intervals
count, bin_edges = np.histogram(df_can['2013'])

df_can['2013'].plot(kind='hist', figsize=(8, 5), xticks=bin_edges)

plt.title('Histogram of Immigration from 195 countries in 2013') # add
a title to the histogram
plt.ylabel('Number of Countries') # add y-label
plt.xlabel('Number of Immigrants') # add x-label

plt.show()
```



Sweden	281	308	222	176	128	158	187	198	171	
182 ...										
	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
Country										
Denmark	89	62	101	97	108	81	92	93	94	81
Norway	73	57	53	73	66	75	46	49	53	59
Sweden	129	205	139	193	165	167	159	134	140	140

[3 rows x 34 columns]

generate histogram

df_can.loc[['Denmark', 'Norway', 'Sweden'], years].plot.hist()

<AxesSubplot:ylabel='Frequency'>



That does not look right!

Don't worry, you'll often come across situations like this when creating plots. The solution often lies in how the underlying dataset is structured.

Instead of plotting the population frequency distribution of the population for the 3 countries, *pandas* instead plotted the population frequency distribution for the **years**.

This can be easily fixed by first transposing the dataset, and then plotting as shown below.


```
# transpose dataframe
df_t = df_can.loc[['Denmark', 'Norway', 'Sweden'], years].transpose()
df_t.head(5)
```

Country	Denmark	Norway	Sweden
1980	272	116	281
1981	293	77	308
1982	299	106	222
1983	106	51	176
1984	93	31	128

```
# generate histogram
df_t.plot(kind='hist', figsize=(10, 6))

plt.title('Histogram of Immigration from Denmark, Norway, and Sweden
from 1980 - 2013')
plt.ylabel('Number of Years')
plt.xlabel('Number of Immigrants')

plt.show()
```



Let's make a few modifications to improve the impact and aesthetics of the previous plot:


- increase the bin size to 15 by passing in `bins` parameter;
- set transparency to 60% by passing in `alpha` parameter;
- label the x-axis by passing in `x-label` parameter;
- change the colors of the plots by passing in `color` parameter.

```
# let's get the x-tick values
count, bin_edges = np.histogram(df_t, 15)

# un-stacked histogram
df_t.plot(kind='hist',
          figsize=(10, 6),
          bins=15,
          alpha=0.6,
          xticks=bin_edges,
          color=['coral', 'darkslateblue', 'mediumseagreen']
        )

plt.title('Histogram of Immigration from Denmark, Norway, and Sweden
from 1980 - 2013')
plt.ylabel('Number of Years')
plt.xlabel('Number of Immigrants')

plt.show()
```



Tip: For a full listing of colors available in Matplotlib, run the following code in your python shell:

```
import matplotlib
for name, hex in matplotlib.colors.cnames.items():
    print(name, hex)
```


If we do not want the plots to overlap each other, we can stack them using the `stacked` parameter. Let's also adjust the min and max x-axis labels to remove the extra gap on the edges of the plot. We can pass a tuple (min,max) using the `xlim` parameter, as show below.

```
count, bin_edges = np.histogram(df_t, 15)
xmin = bin_edges[0] - 10 # first bin value is 31.0, adding buffer
of 10 for aesthetic purposes
xmax = bin_edges[-1] + 10 # last bin value is 308.0, adding buffer
of 10 for aesthetic purposes

# stacked Histogram
df_t.plot(kind='hist',
          figsize=(10, 6),
          bins=15,
          xticks=bin_edges,
          color=['coral', 'darkslateblue', 'mediumseagreen'],
          stacked=True,
          xlim=(xmin, xmax)
        )

plt.title('Histogram of Immigration from Denmark, Norway, and Sweden
from 1980 - 2013')
plt.ylabel('Number of Years')
plt.xlabel('Number of Immigrants')

plt.show()
```



Question: Use the scripting layer to display the immigration distribution for Greece, Albania, and Bulgaria for years 1980 - 2013? Use an overlapping plot with 15 bins and a transparency value of 0.35.

```
### type your answer here
# let's quickly view the dataset
df_can.loc[['Greece', 'Albania', 'Bulgaria'], years]
```


	1980	1981	1982	1983	1984	1985	1986	1987	1988
1989 ... \									
Country									
...									
Greece	1065	953	897	633	580	584	547	765	564
773 ...									
Albania	1	0	0	0	0	0	1	2	2
3 ...									
Bulgaria	24	20	12	33	11	24	33	52	43
85 ...									

	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
Country										
Greece	120	100	74	110	107	119	101	102	146	298
Albania	1450	1223	856	702	560	716	561	539	620	603
Bulgaria	2022	1738	1419	1172	994	784	556	365	451	512

[3 rows x 34 columns]

```
# generate histogram
df_can.loc[['Greece', 'Albania', 'Bulgaria'], years].plot.hist()

<AxesSubplot:ylabel='Frequency'>
```



```

# transpose dataframe
df_t = df_can.loc[['Greece', 'Albania', 'Bulgaria'],
years].transpose()
df_t.head(5)


Country  Greece  Albania  Bulgaria
1980      1065         1         24
1981       953         0         20
1982       897         0         12
1983       633         0         33
1984       580         0         11

# generate histogram
df_t.plot(kind='hist', figsize=(10, 6))

plt.title('Histogram of Immigration from Greece, Albania and Bulgaria
from 1980 - 2013')
plt.ylabel('Number of Years')
plt.xlabel('Number of Immigrants')

plt.show()

```



```

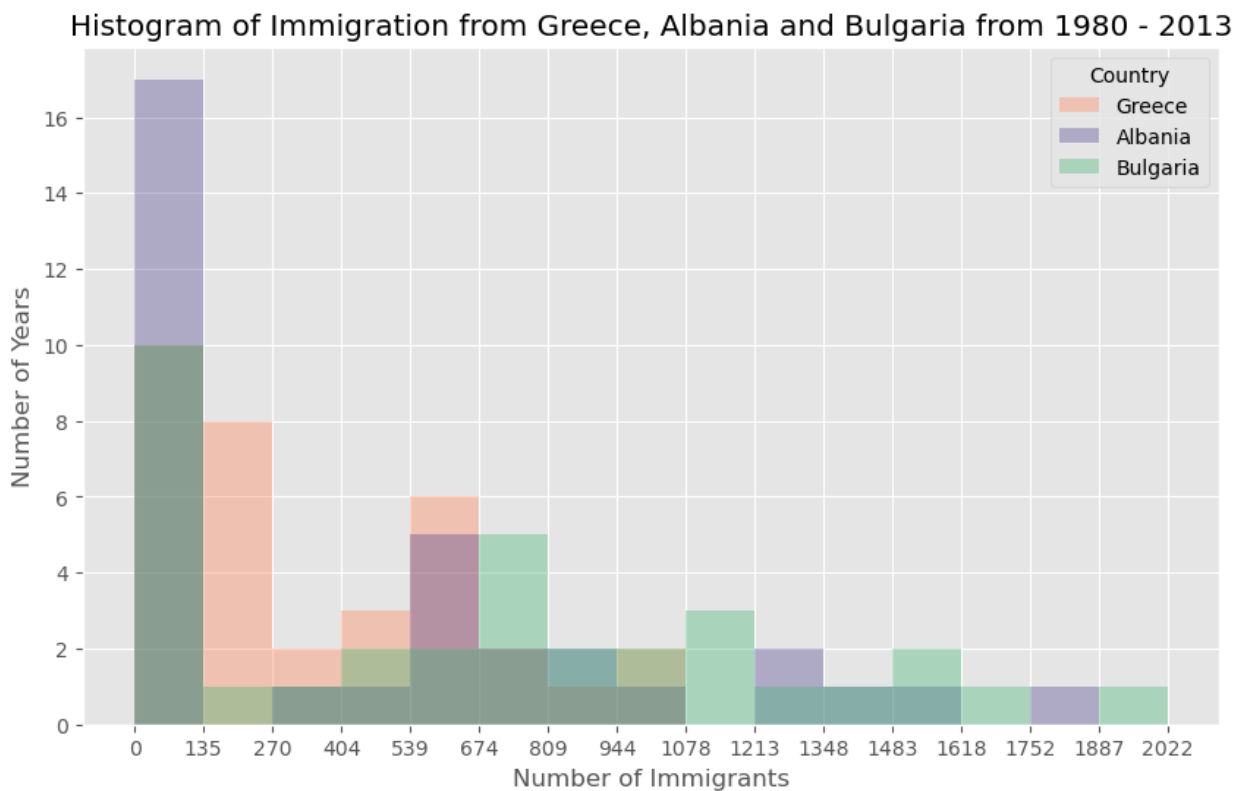
# let's get the x-tick values
count, bin_edges = np.histogram(df_t, 15)

```

```
# un-stacked histogram
df_t.plot(kind='hist',
          figsize=(10, 6),
          bins=15,
          alpha=0.35,
          xticks=bin_edges,
          color=['coral', 'darkslateblue', 'mediumseagreen'])

plt.title('Histogram of Immigration from Greece, Albania and Bulgaria
from 1980 - 2013')
plt.ylabel('Number of Years')
plt.xlabel('Number of Immigrants')

plt.show()
```



Bar Charts (Dataframe)

A bar plot is a way of representing data where the *length* of the bars represents the magnitude/size of the feature/variable. Bar graphs usually represent numerical and categorical variables grouped in intervals.

To create a bar plot, we can pass one of two arguments via `kind` parameter in `plot()`:

- `kind=bar` creates a *vertical* bar plot
- `kind=barh` creates a *horizontal* bar plot

Vertical bar plot

In vertical bar graphs, the x-axis is used for labelling, and the length of bars on the y-axis corresponds to the magnitude of the variable being measured. Vertical bar graphs are particularly useful in analyzing time series data. One disadvantage is that they lack space for text labelling at the foot of each bar.

Let's start off by analyzing the effect of Iceland's Financial Crisis:

The 2008 - 2011 Icelandic Financial Crisis was a major economic and political event in Iceland. Relative to the size of its economy, Iceland's systemic banking collapse was the largest experienced by any country in economic history. The crisis led to a severe economic depression in 2008 - 2011 and significant political unrest.

Question: Let's compare the number of Icelandic immigrants (country = 'Iceland') to Canada from year 1980 to 2013.


```
# step 1: get the data
df_iceland = df_can.loc['Iceland', years]
df_iceland.head()

1980    17
1981    33
1982    10
1983     9
1984    13
Name: Iceland, dtype: object

# step 2: plot data
df_iceland.plot(kind='bar', figsize=(10, 6))

plt.xlabel('Year') # add to x-label to the plot
plt.ylabel('Number of immigrants') # add y-label to the plot
plt.title('Icelandic immigrants to Canada from 1980 to 2013') # add
title to the plot

plt.show()
```



The bar plot above shows the total number of immigrants broken down by each year. We can clearly see the impact of the financial crisis; the number of immigrants to Canada started increasing rapidly after 2008.

Let's annotate this on the plot using the `annotate` method of the **scripting layer** or the **pyplot interface**. We will pass in the following parameters:

- `s`: str, the text of annotation.
- `xy`: Tuple specifying the (x,y) point to annotate (in this case, end point of arrow).
- `xytext`: Tuple specifying the (x,y) point to place the text (in this case, start point of arrow).
- `xycoords`: The coordinate system that xy is given in - 'data' uses the coordinate system of the object being annotated (default).
- `arrowprops`: Takes a dictionary of properties to draw the arrow:
 - `arrowstyle`: Specifies the arrow style, '`->`' is standard arrow.
 - `connectionstyle`: Specifies the connection type. `arc3` is a straight line.
 - `color`: Specifies color of arrow.
 - `lw`: Specifies the line width.

I encourage you to read the Matplotlib documentation for more details on annotations:
https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.annotate.html.


```


df_iceland.plot(kind='bar', figsize=(10, 6), rot=90) # rotate the
xticks(labelled points on x-axis) by 90 degrees

plt.xlabel('Year')
plt.ylabel('Number of Immigrants')
plt.title('Icelandic Immigrants to Canada from 1980 to 2013')

# Annotate arrow
plt.annotate('', # s: str. Will leave it blank for no text
             xy=(32, 70), # place head of the arrow at point (year
2012 , pop 70)
             xytext=(28, 20), # place base of the arrow at point
(year 2008 , pop 20)
             xycoords='data', # will use the coordinate system of the
object being annotated
             arrowprops=dict(arrowstyle='->', connectionstyle='arc3',
color='blue', lw=2)
             )

plt.show()

```



Let's also annotate a text to go over the arrow. We will pass in the following additional parameters:

- **rotation**: rotation angle of text in degrees (counter clockwise)

- `va`: vertical alignment of text ['center' | 'top' | 'bottom' | 'baseline']
- `ha`: horizontal alignment of text ['center' | 'right' | 'left']


```
df_iceland.plot(kind='bar', figsize=(10, 6), rot=90)

plt.xlabel('Year')
plt.ylabel('Number of Immigrants')
plt.title('Icelandic Immigrants to Canada from 1980 to 2013')

# Annotate arrow
plt.annotate('', # s: str. will leave it blank for no text
             xy=(32, 70), # place head of the arrow at point (year
2012 , pop 70)
             xytext=(28, 20), # place base of the arrow at point
(year 2008 , pop 20)
             xycoords='data', # will use the coordinate system of the
object being annotated
             arrowprops=dict(arrowstyle='->', connectionstyle='arc3',
color='blue', lw=2)
             )

# Annotate Text
plt.annotate('2008 - 2011 Financial Crisis', # text to display
             xy=(28, 30), # start the text at at point (year 2008 ,
pop 30)
             rotation=72.5, # based on trial and error to match the
arrow
             va='bottom', # want the text to be vertically 'bottom'
aligned
             ha='left', # want the text to be horizontally 'left'
aligned.
             )

plt.show()
```



Horizontal Bar Plot

Sometimes it is more practical to represent the data horizontally, especially if you need more room for labelling the bars. In horizontal bar graphs, the y-axis is used for labelling, and the length of bars on the x-axis corresponds to the magnitude of the variable being measured. As you will see, there is more room on the y-axis to label categorical variables.

Question: Using the scripting later and the `df_can` dataset, create a *horizontal* bar plot showing the *total* number of immigrants to Canada from the top 15 countries, for the period 1980 - 2013. Label each country with the total immigrant count.

Step 1: Get the data pertaining to the top 15 countries.

```
# sort dataframe on 'Total' column (descending)
df_can.sort_values(by='Total', ascending=True, inplace=True)
```

```
# get top 15 countries
df_top15 = df_can['Total'].tail(15)
df_top15
```

Country	
Romania	93585
Viet Nam	97146
Jamaica	106431
France	109091

Lebanon	115359
Poland	139241
Republic of Korea	142581
Sri Lanka	148358
Iran (Islamic Republic of)	175923
United States of America	241122
Pakistan	241600
Philippines	511391
United Kingdom of Great Britain and Northern Ireland	551500
China	659962
India	691904

Name: Total, dtype: int64

Step 2: Plot data:


1. Use `kind='barh'` to generate a bar chart with horizontal bars.
2. Make sure to choose a good size for the plot and to label your axes and to give the plot a title.
3. Loop through the countries and annotate the immigrant population using the `annotate` function of the scripting interface.

```
# generate plot
df_top15.plot(kind='barh', figsize=(12, 12), color='steelblue')
plt.xlabel('Number of Immigrants')
plt.title('Top 15 Countries Contributing to the Immigration to Canada
between 1980 - 2013')

# annotate value labels to each country
for index, value in enumerate(df_top15):
    label = format(int(value), ',') # format int with commas

# place text at the end of bar (subtracting 47000 from x, and 0.1 from
y to make it fit within the bar)
plt.annotate(label, xy=(value - 47000, index - 0.10), color='white')

plt.show()
```



Thank you for completing this lab!

Author

Alex Aklson

Other Contributors

Jay Rajasekharan, Ehsan M. Kermani, Slobodan Markovic, Weiqing Wang, Pooja.

Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2023-06-11	2.5	Pooja	Changed to work on clean data, links updated
2021-05-29	2.4	Weiqing Wang	Fixed typos and code spells.
2021-01-20	2.3	Lakshmi Holla	Changed TOC cell markdown
2021-01-05	2.2	Lakshmi Holla	Changed solution code for annotate
2020-11-03	2.1	Lakshmi Holla	Changed the URL of excel file

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-08-27	2.0	Lavanya	Moved lab to course repo in GitLab

© IBM Corporation 2020. All rights reserved.