

Exploring and pre-processing a dataset using Pandas

Estimated time needed: **30** minutes

Objectives

After completing this lab you will be able to:

- Explore the dataset
- Pre-process dataset as required (may be for visualization)

Introduction

The aim of this lab is to provide you a refresher on the **Pandas** library, so that you can pre-process and analyse the datasets before applying data visualization techniques on it. This lab will work as a crash course on *pandas*. If you are interested in learning more about the *pandas* library, detailed description and explanation of how to use it and how to clean, munge, and process data stored in a *pandas* dataframe are provided in our course [Data Analysis with Python](#) and [Python for Applied Data Science](#)

Table of Contents

Exploring Datasets with *pandas*

pandas is an essential data analysis toolkit for Python. From their [website](#):

pandas is a Python package providing fast, flexible, and expressive data structures designed to make working with “relational” or “labeled” data both easy and intuitive. It aims to be the fundamental high-level building block for doing practical, **real world** data analysis in Python.

The course heavily relies on *pandas* for data wrangling, analysis, and visualization. We encourage you to spend some time and familiarize yourself with the *pandas* API Reference: <http://pandas.pydata.org/pandas-docs/stable/api.html>.

The Dataset: Immigration to Canada from 1980 to 2013

Dataset Source: [International migration flows to and from selected countries - The 2015 revision](#).

In this lab, we will focus on the Canadian immigration data.

The Canada Immigration dataset can be fetched from here.

The first thing we'll do is install **openpyxl** (formerly **xlrd**), a module that *pandas* requires to read Excel files.



mamba (1.4.2) supported by @QuantStack

GitHub: <https://github.com/mamba-org/mamba>

Twitter: <https://twitter.com/QuantStack>

Looking for: ['openpyxl==3.0.9']

```
ain/linux-64 0.0 B / ???.?MB @ ???.?MB/s
0.1s
pkgs/main/noarch 0.0 B / ???.?MB @ ???.?
MB/s 0.1s
pkgs/r/linux-64 0.0 B / ???.?MB @ ???.?
MB/s 0.1s
pkgs/r/noarch 0.0 B / ???.?MB @ ???.?
MB/s 0.1sain/linux-64 16.4kB / ???.?MB @
106.7kB/s 0.2s
pkgs/main/noarch 41.0kB / ???.?MB @
266.6kB/s 0.2s
pkgs/r/linux-64 45.1kB / ???.?MB @
292.7kB/s 0.2s
pkgs/r/noarch 12.3kB / ???.?MB @
79.9kB/s 0.2sain/linux-64 323.6kB / ???.?MB
@ 1.3MB/s 0.3s
pkgs/main/noarch 524.3kB / ???.?MB @
2.1MB/s 0.3s
pkgs/r/linux-64 507.9kB / ???.?MB @
2.0MB/s 0.3s
pkgs/r/noarch 475.1kB / ???.?MB @
1.9MB/s 0.3sain/noarch 873.1kB @
2.7MB/s 0.3s
[+] 0.4s
pkgs/main/linux-64 872.5kB / ???.?MB @
2.4MB/s 0.4s
pkgs/r/linux-64 987.1kB / ???.?MB @
2.8MB/s 0.4s
pkgs/r/noarch 921.6kB / ???.?MB @
2.6MB/s 0.4sain/linux-64 1.3MB / ???.?MB
@ 2.9MB/s 0.5s
pkgs/r/linux-64 1.5MB / ???.?MB @
3.2MB/s 0.5s
pkgs/r/noarch 1.4MB / ???.?MB @
3.0MB/s 0.5sain/linux-64 1.8MB / ???.?MB
@ 3.2MB/s 0.6s
pkgs/r/noarch 1.9MB / ???.?MB @
```

3.3MB/s 0.6sain/linux-64 2.3MB @ 3.4MB/s
0.7s
pkgs/r/noarch 2.3MB @ 3.5MB/s
Finalizing 0.7sain/linux-64 2.8MB / ???
MB @ 3.6MB/s 0.8sain/linux-64
3.4MB / ???MB @ 3.9MB/s 0.9sain/linux-64
3.8MB / ???MB @ 3.9MB/s
1.0sain/linux-64 4.3MB / ???MB @
3.9MB/s 1.1sain/linux-64 4.8MB / ???MB
@ 4.1MB/s 1.2sain/linux-64
5.3MB / ???MB @ 4.1MB/s 1.3sain/linux-64
5.8MB / ???MB @ 4.2MB/s
1.4sain/linux-64 6.3MB / ???MB @
4.2MB/s 1.5sain/linux-64 6.8MB / ???MB
@ 4.3MB/s 1.6sain/linux-64 7.0MB @
4.3MB/s Finalizing 1.7sain/linux-64 @
4.3MB/s 1.7s
e/jupyterlab/conda/envs/python

Updating specs:

- openpyxl==3.0.9
- ca-certificates
- certifi
- openssl

Package	Version	Build	Channel
Size			

Install:

+ et_xmlfile	1.1.0	py37h06a4308_0	pkgs/main/linux-64
10kB			
+ openpyxl	3.0.9	pyhd3eb1b0_0	pkgs/main/noarch
168kB			

Upgrade:

- ca-certificates	2023.5.7	hbcca054_0	conda-forge
+ ca-certificates	2024.3.11	h06a4308_0	pkgs/main/linux-64
130kB			
- openssl	1.1.1t	h0b41bf4_0	conda-forge

```
+ openssl          1.1.1w  h7f8727e_0      pkgs/main/linux-64
4MB
```

Summary:

Install: 2 packages

Upgrade: 2 packages

Total download: 4MB

```

0.0 B ca-certificates 0.0s
Extracting 0
0.0s 0.0 B ca-certificates 0.1s
Extracting 0
0.0s file 10.2kB @ 61.9kB/s
0.2s
[+] 0.2s
Downloading 4.2MB
0.2s
Extracting (4) 0 ca-certificates
0.0s 0 ca-certificates
0.1s 0 ca-certificates
0.2s 0 ca-certificates
0.3s 0 et_xmlfile
0.4s 0 et_xmlfile
0.5s 0 et_xmlfile
0.6s 0 et_xmlfile
0.7s 0 openpyxl
0.8s 0 openpyxl
0.9s 0 openpyxl
1.0s 0 openpyxl
1.1s 0 openssl
1.2s 0 openssl
1.3s 0 openssl
1.4s 1 openssl
1.5s 1 ca-certificates 1.6s
3 openssl 1.7s 3 openssl
1.8s
```

Next, we'll do is import two key data analysis modules: *pandas* and *numpy*.

```
import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library
```

Let's download and import our primary Canadian Immigration dataset using *pandas*'s `read_excel()` method.

```
df_can = pd.read_excel(
    'https://cf-courses-data.s3.us.cloud-object-
    storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DV0101EN-
    SkillsNetwork/Data%20Files/Canada.xlsx',
    sheet_name='Canada by Citizenship',
    skiprows=range(20),
    skipfooter=2)

print('Data read into a pandas dataframe!')
```

Data read into a pandas dataframe!

Let's view the top 5 rows of the dataset using the `head()` function.

```
df_can.head()
# tip: You can specify the number of rows you'd like to see as
# follows: df_can.head(10)
```

	Type	Coverage	OdName	AREA	AreaName	REG	\
0	Immigrants	Foreigners	Afghanistan	935	Asia	5501	
1	Immigrants	Foreigners	Albania	908	Europe	925	
2	Immigrants	Foreigners	Algeria	903	Africa	912	
3	Immigrants	Foreigners	American Samoa	909	Oceania	957	
4	Immigrants	Foreigners	Andorra	908	Europe	925	

	RegName	DEV	DevName	1980	...	2004	2005
2006	\						
0	Southern Asia	902	Developing regions	16	...	2978	3436
3009							
1	Southern Europe	901	Developed regions	1	...	1450	1223
856							
2	Northern Africa	902	Developing regions	80	...	3616	3626
4807							
3	Polynesia	902	Developing regions	0	...	0	0
1							
4	Southern Europe	901	Developed regions	0	...	0	0
1							

	2007	2008	2009	2010	2011	2012	2013
0	2652	2111	1746	1758	2203	2635	2004
1	702	560	716	561	539	620	603
2	3623	4005	5393	4752	4325	3774	4331
3	0	0	0	0	0	0	0
4	1	0	0	0	0	1	1

[5 rows x 43 columns]

We can also view the bottom 5 rows of the dataset using the `tail()` function.

```
df_can.tail()
```

	Type	Coverage	OdName	AREA	AreaName	REG	\
190	Immigrants	Foreigners	Viet Nam	935	Asia	920	
191	Immigrants	Foreigners	Western Sahara	903	Africa	912	
192	Immigrants	Foreigners	Yemen	935	Asia	922	
193	Immigrants	Foreigners	Zambia	903	Africa	910	
194	Immigrants	Foreigners	Zimbabwe	903	Africa	910	

	RegName	DEV	DevName	1980	...	2004
2005	2006	\				
190	South-Eastern Asia	902	Developing regions	1191	...	1816
1852	3153					
191	Northern Africa	902	Developing regions	0	...	0
0	1					
192	Western Asia	902	Developing regions	1	...	124
161	140					
193	Eastern Africa	902	Developing regions	11	...	56
91	77					
194	Eastern Africa	902	Developing regions	72	...	1450
615	454					

	2007	2008	2009	2010	2011	2012	2013
190	2574	1784	2171	1942	1723	1731	2112
191	0	0	0	0	0	0	0
192	122	133	128	211	160	174	217
193	71	64	60	102	69	46	59
194	663	611	508	494	434	437	407

```
[5 rows x 43 columns]
```

When analyzing a dataset, it's always a good idea to start by getting basic information about your dataframe. We can do this by using the `info()` method.

This method can be used to get a short summary of the dataframe.

```
df_can.info(verbose=False)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 195 entries, 0 to 194
Columns: 43 entries, Type to 2013
dtypes: int64(37), object(6)
memory usage: 65.6+ KB
```

To get the list of column headers we can call upon the data frame's `columns` instance variable.

```
df_can.columns
```

```
Index(['Type', 'Coverage', 'OdName', 'AREA', 'AreaName', 'REG',
```

```

1982,      'RegName',      'DEV',  'DevName',      1980,      1981,
1983,      1983,      1984,      1985,      1986,      1987,
1988,      1989,      1990,      1991,      1992,      1993,
1994,      1995,      1996,      1997,      1998,      1999,
2000,      2001,      2002,      2003,      2004,      2005,
2006,      2007,      2008,      2009,      2010,      2011,
2012,      2013],
dtype='object')

```

Similarly, to get the list of indices we use the `.index` instance variables.

```

df_can.index
RangeIndex(start=0, stop=195, step=1)

```

Note: The default type of instance variables `index` and `columns` are **NOT** list.

```

print(type(df_can.columns))
print(type(df_can.index))

<class 'pandas.core.indexes.base.Index'>
<class 'pandas.core.indexes.range.RangeIndex'>

```

To get the index and columns as lists, we can use the `tolist()` method.

```

df_can.columns.tolist()

['Type',
 'Coverage',
 'OdName',
 'AREA',
 'AreaName',
 'REG',
 'RegName',
 'DEV',
 'DevName',
 1980,
 1981,
 1982,
 1983,
 1984,
 1985,
 1986,

```



```
1987,  
1988,  
1989,  
1990,  
1991,  
1992,  
1993,  
1994,  
1995,  
1996,  
1997,  
1998,  
1999,  
2000,  
2001,  
2002,  
2003,  
2004,  
2005,  
2006,  
2007,  
2008,  
2009,  
2010,  
2011,  
2012,  
2013]
```

```
df_can.index.tolist()
```

```
[0,  
1,  
2,  
3,  
4,  
5,  
6,  
7,  
8,  
9,  
10,  
11,  
12,  
13,  
14,  
15,  
16,  
17,  
18,  
19,
```

20,
21,
22,
23,
24,
25,
26,
27,
28,
29,
30,
31,
32,
33,
34,
35,
36,
37,
38,
39,
40,
41,
42,
43,
44,
45,
46,
47,
48,
49,
50,
51,
52,
53,
54,
55,
56,
57,
58,
59,
60,
61,
62,
63,
64,
65,
66,
67,
68,

69,
70,
71,
72,
73,
74,
75,
76,
77,
78,
79,
80,
81,
82,
83,
84,
85,
86,
87,
88,
89,
90,
91,
92,
93,
94,
95,
96,
97,
98,
99,
100,
101,
102,
103,
104,
105,
106,
107,
108,
109,
110,
111,
112,
113,
114,
115,
116,
117,

118,
119,
120,
121,
122,
123,
124,
125,
126,
127,
128,
129,
130,
131,
132,
133,
134,
135,
136,
137,
138,
139,
140,
141,
142,
143,
144,
145,
146,
147,
148,
149,
150,
151,
152,
153,
154,
155,
156,
157,
158,
159,
160,
161,
162,
163,
164,
165,
166,

```
167,  
168,  
169,  
170,  
171,  
172,  
173,  
174,  
175,  
176,  
177,  
178,  
179,  
180,  
181,  
182,  
183,  
184,  
185,  
186,  
187,  
188,  
189,  
190,  
191,  
192,  
193,  
194]
```

```
print(type(df_can.columns.tolist()))  
print(type(df_can.index.tolist()))  
  
<class 'list'>  
<class 'list'>
```

To view the dimensions of the dataframe, we use the `shape` instance variable of it.

```
# size of dataframe (rows, columns)  
df_can.shape  
  
(195, 43)
```

Note: The main types stored in *pandas* objects are `float`, `int`, `bool`, `datetime64[ns]`, `datetime64[ns, tz]`, `timedelta[ns]`, `category`, and `object` (string). In addition, these dtypes have item sizes, e.g. `int64` and `int32`.

Let's clean the data set to remove a few unnecessary columns. We can use *pandas* `drop()` method as follows:

```
# in pandas axis=0 represents rows (default) and axis=1 represents columns.
```

```
df_can.drop(['AREA', 'REG', 'DEV', 'Type', 'Coverage'], axis=1,
inplace=True)
df_can.head(2)
```

	OdName	AreaName	RegName	DevName	1980
1981 \					
0	Afghanistan	Asia	Southern Asia	Developing regions	16
39					
1	Albania	Europe	Southern Europe	Developed regions	1
0					

	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009
2010 \											
0	39	47	71	340	...	2978	3436	3009	2652	2111	1746
1758											
1	0	0	0	0	...	1450	1223	856	702	560	716
561											

	2011	2012	2013
0	2203	2635	2004
1	539	620	603

[2 rows x 38 columns]

Let's rename the columns so that they make sense. We can use `rename()` method by passing in a dictionary of old and new names as follows:

```
df_can.rename(columns={'OdName': 'Country', 'AreaName': 'Continent',
'RegName': 'Region'}, inplace=True)
df_can.columns
```

```
Index([ 'Country', 'Continent', 'Region', 'DevName',
1980,
1981, 1982, 1983, 1984,
1985,
1986, 1987, 1988, 1989,
1990,
1991, 1992, 1993, 1994,
1995,
1996, 1997, 1998, 1999,
2000,
2001, 2002, 2003, 2004,
2005,
2006, 2007, 2008, 2009,
2010,
2011, 2012, 2013],
dtype='object')
```

We will also add a 'Total' column that sums up the total immigrants by country over the entire period 1980 - 2013, as follows:

```
df_can['Total'] = df_can.sum(axis=1)
df_can['Total']

/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages/
ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns
in DataFrame reductions (with 'numeric_only=None') is deprecated; in a
future version this will raise TypeError. Select only valid columns
before calling the reduction.
    """Entry point for launching an IPython kernel.
```

0	58639
1	15699
2	69439
3	6
4	15
...	
190	97146
191	2
192	2985
193	1677
194	8598

Name: Total, Length: 195, dtype: int64

We can check to see how many null objects we have in the dataset as follows:

```
df_can.isnull().sum()
```

Country	0
Continent	0
Region	0
DevName	0
1980	0
1981	0
1982	0
1983	0
1984	0
1985	0
1986	0
1987	0
1988	0
1989	0
1990	0
1991	0
1992	0
1993	0
1994	0
1995	0

```

1996      0
1997      0
1998      0
1999      0
2000      0
2001      0
2002      0
2003      0
2004      0
2005      0
2006      0
2007      0
2008      0
2009      0
2010      0
2011      0
2012      0
2013      0
Total      0
dtype: int64

```

Finally, let's view a quick summary of each column in our dataframe using the `describe()` method.

```

df_can.describe()

```

	1980	1981	1982	1983
count	195.000000	195.000000	195.000000	195.000000
mean	508.394872	566.989744	534.723077	387.435897
std	1949.588546	2152.643752	1866.997511	1204.333597
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	13.000000	10.000000	11.000000	12.000000
75%	251.500000	295.500000	275.000000	173.000000
max	22045.000000	24796.000000	20620.000000	10015.000000

	1985	1986	1987	1988
count	195.000000	195.000000	195.000000	195.000000

mean	358.861538	441.271795	691.133333	714.389744
843.241026				
std	1079.309600	1225.576630	2109.205607	2443.606788
2555.048874				
min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	0.000000	0.500000	0.500000	1.000000
1.000000				
50%	17.000000	18.000000	26.000000	34.000000
44.000000				
75%	197.000000	254.000000	434.000000	409.000000
508.500000				
max	9564.000000	9470.000000	21337.000000	27359.000000
23795.000000				

	...	2005	2006	2007	2008	\
count	...	195.000000	195.000000	195.000000	195.000000	
mean	...	1320.292308	1266.958974	1191.820513	1246.394872	
std	...	4425.957828	3926.717747	3443.542409	3694.573544	
min	...	0.000000	0.000000	0.000000	0.000000	
25%	...	28.500000	25.000000	31.000000	31.000000	
50%	...	210.000000	218.000000	198.000000	205.000000	
75%	...	832.000000	842.000000	899.000000	934.500000	
max	...	42584.000000	33848.000000	28742.000000	30037.000000	

		2009	2010	2011	2012
2013	\				
count		195.000000	195.000000	195.000000	195.000000
195.000000					
mean		1275.733333	1420.287179	1262.533333	1313.958974
1320.702564					
std		3829.630424	4462.946328	4030.084313	4247.555161
4237.951988					
min		0.000000	0.000000	0.000000	0.000000
0.000000					
25%		36.000000	40.500000	37.500000	42.500000
45.000000					
50%		214.000000	211.000000	179.000000	233.000000
213.000000					
75%		888.000000	932.000000	772.000000	783.000000
796.000000					
max		29622.000000	38617.000000	36765.000000	34315.000000
34129.000000					

	Total
count	195.000000
mean	32867.451282
std	91785.498686
min	1.000000
25%	952.000000

```
50%      5018.000000
75%     22239.500000
max      691904.000000
```

```
[8 rows x 35 columns]
```

pandas Intermediate: Indexing and Selection (slicing)

Select Column

There are two ways to filter on a column name:

Method 1: Quick and easy, but only works if the column name does NOT have spaces or special characters.

```
df.column_name          # returns series
```

Method 2: More robust, and can filter on multiple columns.

```
df['column']             # returns series
df[['column 1', 'column 2']] # returns dataframe
```

Example: Let's try filtering on the list of countries ('Country').

```
df_can.Country          # returns a series
0      Afghanistan
1      Albania
2      Algeria
3      American Samoa
4      Andorra
...
190     Viet Nam
191  Western Sahara
192      Yemen
193      Zambia
194      Zimbabwe
Name: Country, Length: 195, dtype: object
```

Let's try filtering on the list of countries ('Country') and the data for years: 1980 - 1985.

```
df_can[['Country', 1980, 1981, 1982, 1983, 1984, 1985]] # returns a
dataframe
```

```
# notice that 'Country' is string, and the years are integers.
# for the sake of consistency, we will convert all column names to
string later on.
```

	Country	1980	1981	1982	1983	1984	1985
0	Afghanistan	16	39	39	47	71	340
1	Albania	1	0	0	0	0	0
2	Algeria	80	67	71	69	63	44
3	American Samoa	0	1	0	0	0	0
4	Andorra	0	0	0	0	0	0
...
190	Viet Nam	1191	1829	2162	3404	7583	5907
191	Western Sahara	0	0	0	0	0	0
192	Yemen	1	2	1	6	0	18
193	Zambia	11	17	11	7	16	9
194	Zimbabwe	72	114	102	44	32	29

[195 rows x 7 columns]

Select Row

There are main 2 ways to select rows:

```
df.loc[label]    # filters by the labels of the index/column
df.iloc[index]   # filters by the positions of the index/column
```

Before we proceed, notice that the default index of the dataset is a numeric range from 0 to 194. This makes it very difficult to do a query by a specific country. For example to search for data on Japan, we need to know the corresponding index value.

This can be fixed very easily by setting the 'Country' column as the index using `set_index()` method.

```
df_can.set_index('Country', inplace=True)
# tip: The opposite of set is reset. So to reset the index, we can use
df_can.reset_index()
```

```
df_can.head(3)
```

	Continent	Region	DevName	1980	1981
1982 \ Country					
Afghanistan	Asia	Southern Asia	Developing regions	16	39
Albania	Europe	Southern Europe	Developed regions	1	0
Algeria	Africa	Northern Africa	Developing regions	80	67

	1983	1984	1985	1986	...	2005	2006	2007	2008	2009
2010 \ Country					...					
Afghanistan	47	71	340	496	...	3436	3009	2652	2111	1746
1758										
Albania	0	0	0	1	...	1223	856	702	560	716
561										
Algeria	69	63	44	69	...	3626	4807	3623	4005	5393
4752										
	2011	2012	2013	Total						
Country										
Afghanistan	2203	2635	2004	58639						
Albania	539	620	603	15699						
Algeria	4325	3774	4331	69439						

[3 rows x 38 columns]

optional: to remove the name of the index
df_can.index.name = None

Example: Let's view the number of immigrants from Japan (row 87) for the following scenarios:
1. The full row data (all columns) 2. For year 2013 3. For years 1980 to 1985

1. the full row data (all columns)
df_can.loc['Japan']

Continent	Asia
Region	Eastern Asia
DevName	Developed regions
1980	701
1981	756
1982	598
1983	309
1984	246
1985	198
1986	248
1987	422
1988	324
1989	494
1990	379
1991	506
1992	605
1993	907
1994	956
1995	826
1996	994
1997	924
1998	897

1999	1083
2000	1010
2001	1092
2002	806
2003	817
2004	973
2005	1067
2006	1212
2007	1250
2008	1284
2009	1194
2010	1168
2011	1265
2012	1214
2013	982
Total	27707

Name: Japan, dtype: object

```
# alternate methods
df_can.iloc[87]
```

Continent	Asia
Region	Eastern Asia
DevName	Developed regions
1980	701
1981	756
1982	598
1983	309
1984	246
1985	198
1986	248
1987	422
1988	324
1989	494
1990	379
1991	506
1992	605
1993	907
1994	956
1995	826
1996	994
1997	924
1998	897
1999	1083
2000	1010
2001	1092
2002	806
2003	817
2004	973
2005	1067

2006	1212
2007	1250
2008	1284
2009	1194
2010	1168
2011	1265
2012	1214
2013	982
Total	27707

Name: Japan, dtype: object

```
df_can[df_can.index == 'Japan']
```

	Continent	Region	DevName	1980	1981	1982
1983	\					
Japan	Asia	Eastern Asia	Developed regions	701	756	598
309						

	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011
2012	\										
Japan	246	198	248	...	1067	1212	1250	1284	1194	1168	1265
1214											

	2013	Total
Japan	982	27707

[1 rows x 38 columns]

2. for year 2013

```
df_can.loc['Japan', 2013]
```

982

alternate method

year 2013 is the last column, with a positional index of 36

```
df_can.iloc[87, 36]
```

982

3. for years 1980 to 1985

```
df_can.loc['Japan', [1980, 1981, 1982, 1983, 1984, 1984]]
```

1980	701
1981	756
1982	598
1983	309
1984	246
1984	246

Name: Japan, dtype: object

Alternative Method

```
df_can.iloc[87, [3, 4, 5, 6, 7, 8]]
```

```
1980    701
1981    756
1982    598
1983    309
1984    246
1985    198
Name: Japan, dtype: object
```

Exercise: Let's view the number of immigrants from **Haiti** for the following scenarios: 1. The full row data (all columns) 2. For year 2000 3. For years 1990 to 1995

```
df_can.loc['Haiti']
df_can.loc['Haiti', 2000]
df_can.loc['Haiti', [1990, 1991, 1992, 1993, 1994, 1995]]

1990    2379
1991    2829
1992    2399
1993    3655
1994    2100
1995    2014
Name: Haiti, dtype: object
```

Column names that are integers (such as the years) might introduce some confusion. For example, when we are referencing the year 2013, one might confuse that when the 2013th positional index.

To avoid this ambiguity, let's convert the column names into strings: '1980' to '2013'.

```
df_can.columns = list(map(str, df_can.columns))
# [print (type(x)) for x in df_can.columns.values] #<-- uncomment to
check type of column headers
```

Since we converted the years to string, let's declare a variable that will allow us to easily call upon the full range of years:

```
# useful for plotting later on
years = list(map(str, range(1980, 2014)))
years

['1980',
 '1981',
 '1982',
 '1983',
 '1984',
 '1985',
 '1986',
 '1987',
 '1988',
```

```
'1989',  
'1990',  
'1991',  
'1992',  
'1993',  
'1994',  
'1995',  
'1996',  
'1997',  
'1998',  
'1999',  
'2000',  
'2001',  
'2002',  
'2003',  
'2004',  
'2005',  
'2006',  
'2007',  
'2008',  
'2009',  
'2010',  
'2011',  
'2012',  
'2013']
```

Exercise: Create a list named 'year' using map function for years ranging from 1990 to 2013. Then extract the data series from the dataframe df_can for Haiti using year list.

```
year = list(map(str, range(1990, 2014)))  
haiti = df_can.loc['Haiti', year]
```

Filtering based on a criteria

To filter the dataframe based on a condition, we simply pass the condition as a boolean vector.

For example, Let's filter the dataframe to show the data on Asian countries (AreaName = Asia).

```
# 1. create the condition boolean series  
condition = df_can['Continent'] == 'Asia'  
print(condition)
```

Afghanistan	True
Albania	False
Algeria	False
American Samoa	False
Andorra	False
...	...
Viet Nam	True


```

Western Sahara    False
Yemen             True
Zambia           False
Zimbabwe         False
Name: Continent, Length: 195, dtype: bool

```

```

# 2. pass this condition into the dataframe
df_can[condition]

```

	Continent
Region \	
Afghanistan	Asia
Southern Asia	
Armenia	Asia
Western Asia	
Azerbaijan	Asia
Western Asia	
Bahrain	Asia
Western Asia	
Bangladesh	Asia
Southern Asia	
Bhutan	Asia
Southern Asia	
Brunei Darussalam	Asia South-
Eastern Asia	
Cambodia	Asia South-
Eastern Asia	
China	Asia
Eastern Asia	
China, Hong Kong Special Administrative Region	Asia
Eastern Asia	
China, Macao Special Administrative Region	Asia
Eastern Asia	
Cyprus	Asia
Western Asia	
Democratic People's Republic of Korea	Asia
Eastern Asia	
Georgia	Asia
Western Asia	
India	Asia
Southern Asia	
Indonesia	Asia South-
Eastern Asia	
Iran (Islamic Republic of)	Asia
Southern Asia	
Iraq	Asia
Western Asia	
Israel	Asia
Western Asia	
Japan	Asia

Eastern Asia		
Jordan	Asia	
Western Asia		
Kazakhstan	Asia	
Central Asia		
Kuwait	Asia	
Western Asia		
Kyrgyzstan	Asia	
Central Asia		
Lao People's Democratic Republic	Asia	South-
Eastern Asia		
Lebanon	Asia	
Western Asia		
Malaysia	Asia	South-
Eastern Asia		
Maldives	Asia	
Southern Asia		
Mongolia	Asia	
Eastern Asia		
Myanmar	Asia	South-
Eastern Asia		
Nepal	Asia	
Southern Asia		
Oman	Asia	
Western Asia		
Pakistan	Asia	
Southern Asia		
Philippines	Asia	South-
Eastern Asia		
Qatar	Asia	
Western Asia		
Republic of Korea	Asia	
Eastern Asia		
Saudi Arabia	Asia	
Western Asia		
Singapore	Asia	South-
Eastern Asia		
Sri Lanka	Asia	
Southern Asia		
State of Palestine	Asia	
Western Asia		
Syrian Arab Republic	Asia	
Western Asia		
Tajikistan	Asia	
Central Asia		
Thailand	Asia	South-
Eastern Asia		
Turkey	Asia	
Western Asia		

Turkmenistan	Asia
Central Asia	
United Arab Emirates	Asia
Western Asia	
Uzbekistan	Asia
Central Asia	
Viet Nam	Asia South-
Eastern Asia	
Yemen	Asia
Western Asia	

	DevName
1980 \	
Afghanistan	Developing regions
16	
Armenia	Developing regions
0	
Azerbaijan	Developing regions
0	
Bahrain	Developing regions
0	
Bangladesh	Developing regions
83	
Bhutan	Developing regions
0	
Brunei Darussalam	Developing regions
79	
Cambodia	Developing regions
12	
China	Developing regions
5123	
China, Hong Kong Special Administrative Region	Developing regions
0	
China, Macao Special Administrative Region	Developing regions
0	
Cyprus	Developing regions
132	
Democratic People's Republic of Korea	Developing regions
1	
Georgia	Developing regions
0	
India	Developing regions
8880	
Indonesia	Developing regions
186	
Iran (Islamic Republic of)	Developing regions
1172	
Iraq	Developing regions
262	

Israel	Developing regions
1403	
Japan	Developed regions
701	
Jordan	Developing regions
177	
Kazakhstan	Developing regions
0	
Kuwait	Developing regions
1	
Kyrgyzstan	Developing regions
0	
Lao People's Democratic Republic	Developing regions
11	
Lebanon	Developing regions
1409	
Malaysia	Developing regions
786	
Maldives	Developing regions
0	
Mongolia	Developing regions
0	
Myanmar	Developing regions
80	
Nepal	Developing regions
1	
Oman	Developing regions
0	
Pakistan	Developing regions
978	
Philippines	Developing regions
6051	
Qatar	Developing regions
0	
Republic of Korea	Developing regions
1011	
Saudi Arabia	Developing regions
0	
Singapore	Developing regions
241	
Sri Lanka	Developing regions
185	
State of Palestine	Developing regions
0	
Syrian Arab Republic	Developing regions
315	
Tajikistan	Developing regions
0	
Thailand	Developing regions

56				
Turkey	Developing regions			
481				
Turkmenistan	Developing regions			
0				
United Arab Emirates	Developing regions			
0				
Uzbekistan	Developing regions			
0				
Viet Nam	Developing regions			
1191				
Yemen	Developing regions			
1				
		1981	1982	1983 1984
1985 \				
Afghanistan		39	39	47 71
340				
Armenia		0	0	0 0
0				
Azerbaijan		0	0	0 0
0				
Bahrain		2	1	1 1
3				
Bangladesh		84	86	81 98
92				
Bhutan		0	0	0 1
0				
Brunei Darussalam		6	8	2 2
4				
Cambodia		19	26	33 10
7				
China		6682	3308	1863 1527
1816				
China, Hong Kong Special Administrative Region		0	0	0 0
0				
China, Macao Special Administrative Region		0	0	0 0
0				
Cyprus		128	84	46 46
43				
Democratic People's Republic of Korea		1	3	1 4
3				
Georgia		0	0	0 0
0				
India		8670	8147	7338 5704
4211				
Indonesia		178	252	115 123
100				
Iran (Islamic Republic of)		1429	1822	1592 1977

1648				
Iraq	245	260	380	428
231				
Israel	1711	1334	541	446
680				
Japan	756	598	309	246
198				
Jordan	160	155	113	102
179				
Kazakhstan	0	0	0	0
0				
Kuwait	0	8	2	1
4				
Kyrgyzstan	0	0	0	0
0				
Lao People's Democratic Republic	6	16	16	7
17				
Lebanon	1119	1159	789	1253
1683				
Malaysia	816	813	448	384
374				
Maldives	0	0	1	0
0				
Mongolia	0	0	0	0
0				
Myanmar	62	46	31	41
23				
Nepal	1	6	1	2
4				
Oman	0	0	8	0
0				
Pakistan	972	1201	900	668
514				
Philippines	5921	5249	4562	3801
3150				
Qatar	0	0	0	0
0				
Republic of Korea	1456	1572	1081	847
962				
Saudi Arabia	0	1	4	1
2				
Singapore	301	337	169	128
139				
Sri Lanka	371	290	197	1086
845				
State of Palestine	0	0	0	0
0				
Syrian Arab Republic	419	409	269	264
385				

Tajikistan 0	0	0	0	0
Thailand 66	53	113	65	82
Turkey 202	874	706	280	338
Turkmenistan 0	0	0	0	0
United Arab Emirates 0	2	2	1	2
Uzbekistan 0	0	0	0	0
Viet Nam 5907	1829	2162	3404	7583
Yemen 18	2	1	6	0
	1986	...	2005	
2006 \				
Afghanistan 3009	496	...	3436	
Armenia 218	0	...	224	
Azerbaijan 236	0	...	359	
Bahrain 12	0	...	12	
Bangladesh 4014	486	...	4171	
Bhutan 10	0	...	5	
Brunei Darussalam 5	12	...	4	
Cambodia 529	8	...	370	
China 33518	1960	...	42584	
China, Hong Kong Special Administrative Region 712	0	...	729	
China, Macao Special Administrative Region 32	0	...	21	
Cyprus 9	48	...	7	
Democratic People's Republic of Korea 10	0	...	14	
Georgia 125	0	...	114	
India 33848	7150	...	36210	

Indonesia 613	127	...	632
Iran (Islamic Republic of) 7480	1794	...	5837
Iraq 1788	265	...	2226
Israel 2625	1212	...	2446
Japan 1212	248	...	1067
Jordan 1827	181	...	1940
Kazakhstan 408	0	...	506
Kuwait 35	4	...	66
Kyrgyzstan 161	0	...	173
Lao People's Democratic Republic 74	21	...	42
Lebanon 3802	2576	...	3709
Malaysia 580	425	...	593
Maldives 0	0	...	0
Mongolia 64	0	...	59
Myanmar 953	18	...	210
Nepal 540	13	...	607
Oman 18	0	...	14
Pakistan 13127	691	...	14314
Philippines 18400	4166	...	18139
Qatar 2	1	...	11
Republic of Korea 6215	1208	...	5832
Saudi Arabia 252	5	...	198
Singapore 298	205	...	392
Sri Lanka 4714	1838	...	4930
State of Palestine	0	...	453

627				
Syrian Arab Republic	493	...	1458	
1145				
Tajikistan	0	...	85	
46				
Thailand	78	...	575	
500				
Turkey	257	...	2065	
1638				
Turkmenistan	0	...	40	
26				
United Arab Emirates	5	...	31	
42				
Uzbekistan	0	...	330	
262				
Viet Nam	2741	...	1852	
3153				
Yemen	7	...	161	
140				
	2007	2008	2009	
2010 \				
Afghanistan	2652	2111	1746	
1758				
Armenia	198	205	267	
252				
Azerbaijan	203	125	165	
209				
Bahrain	22	9	35	
28				
Bangladesh	2897	2939	2104	
4721				
Bhutan	7	36	865	
1464				
Brunei Darussalam	11	10	5	
12				
Cambodia	460	354	203	
200				
China	27642	30037	29622	
30391				
China, Hong Kong Special Administrative Region	674	897	657	
623				
China, Macao Special Administrative Region	16	12	21	
21				
Cyprus	4	7	6	
18				
Democratic People's Republic of Korea	7	19	11	
45				
Georgia	132	112	128	

126			
India	28742	28261	29456
34235			
Indonesia	657	661	504
712			
Iran (Islamic Republic of)	6974	6475	6580
7477			
Iraq	2406	3543	5450
5941			
Israel	2401	2562	2316
2755			
Japan	1250	1284	1194
1168			
Jordan	1421	1581	1235
1831			
Kazakhstan	436	394	431
377			
Kuwait	62	53	68
67			
Kyrgyzstan	135	168	173
157			
Lao People's Democratic Republic	53	32	39
54			
Lebanon	3467	3566	3077
3432			
Malaysia	600	658	640
802			
Maldives	2	1	7
4			
Mongolia	82	59	118
169			
Myanmar	1887	975	1153
556			
Nepal	511	581	561
1392			
Oman	16	10	7
14			
Pakistan	10124	8994	7217
6811			
Philippines	19837	24887	28573
38617			
Qatar	5	9	6
18			
Republic of Korea	5920	7294	5874
5537			
Saudi Arabia	188	249	246
330			
Singapore	690	734	366
805			

Sri Lanka 4422	4123	4756	4547
State of Palestine 654	441	481	400
Syrian Arab Republic 1039	1056	919	917
Tajikistan 52	44	15	50
Thailand 499	487	519	512
Turkey 1492	1463	1122	1238
Turkmenistan 30	37	13	20
United Arab Emirates 86	37	33	37
Uzbekistan 289	284	215	288
Viet Nam 1942	2574	1784	2171
Yemen 211	122	133	128
	2011	2012	2013
Total			
Afghanistan 58639	2203	2635	2004
Armenia 3310	236	258	207
Azerbaijan 2649	138	161	57
Bahrain 475	21	39	32
Bangladesh 65568	2694	2640	3789
Bhutan 5876	1879	1075	487
Brunei Darussalam 600	6	3	6
Cambodia 6538	196	233	288
China 659962	28502	33024	34129
China, Hong Kong Special Administrative Region 9327	591	728	774
China, Macao Special Administrative Region 284	13	33	29
Cyprus 1126	6	12	16
Democratic People's Republic of Korea	97	66	17

388			
Georgia	139	147	125
2068			
India	27509	30933	33087
691904			
Indonesia	390	395	387
13150			
Iran (Islamic Republic of)	7479	7534	11291
175923			
Iraq	6196	4041	4918
69789			
Israel	1970	2134	1945
66508			
Japan	1265	1214	982
27707			
Jordan	1635	1206	1255
35406			
Kazakhstan	381	462	348
8490			
Kuwait	58	73	48
2025			
Kyrgyzstan	159	278	123
2353			
Lao People's Democratic Republic	22	25	15
1089			
Lebanon	3072	1614	2172
115359			
Malaysia	409	358	204
24417			
Maldives	3	1	1
30			
Mongolia	103	68	99
952			
Myanmar	368	193	262
9245			
Nepal	1129	1185	1308
10222			
Oman	10	13	11
224			
Pakistan	7468	11227	12603
241600			
Philippines	36765	34315	29544
511391			
Qatar	3	14	6
157			
Republic of Korea	4588	5316	4509
142581			
Saudi Arabia	278	286	267
3425			

Singapore	219	146	141
14579			
Sri Lanka	3309	3338	2394
148358			
State of Palestine	555	533	462
6512			
Syrian Arab Republic	1005	650	1009
31485			
Tajikistan	47	34	39
503			
Thailand	396	296	400
9174			
Turkey	1257	1068	729
31781			
Turkmenistan	20	20	14
310			
United Arab Emirates	60	54	46
836			
Uzbekistan	162	235	167
3368			
Viet Nam	1723	1731	2112
97146			
Yemen	160	174	217
2985			

[49 rows x 38 columns]

we can pass multiple criteria in the same line.

let's filter for AreaName = Asia and RegName = Southern Asia

```
df_can[(df_can['Continent']=='Asia') & (df_can['Region']=='Southern Asia')]
```

note: When using 'and' and 'or' operators, pandas requires we use '&' and '|' instead of 'and' and 'or'

don't forget to enclose the two conditions in parentheses

DevName	1980	\	Continent	Region	
Afghanistan			Asia	Southern Asia	Developing
regions	16				
Bangladesh			Asia	Southern Asia	Developing
regions	83				
Bhutan			Asia	Southern Asia	Developing
regions	0				
India			Asia	Southern Asia	Developing
regions	8880				
Iran (Islamic Republic of)			Asia	Southern Asia	Developing
regions	1172				
Maldives			Asia	Southern Asia	Developing

regions	0							
Nepal		Asia	Southern Asia	Developing				
regions	1							
Pakistan		Asia	Southern Asia	Developing				
regions	978							
Sri Lanka		Asia	Southern Asia	Developing				
regions	185							
		1981	1982	1983	1984	1985	1986	...
2005 \								
Afghanistan		39	39	47	71	340	496	...
3436								
Bangladesh		84	86	81	98	92	486	...
4171								
Bhutan		0	0	0	1	0	0	...
5								
India		8670	8147	7338	5704	4211	7150	...
36210								
Iran (Islamic Republic of)		1429	1822	1592	1977	1648	1794	...
5837								
Maldives		0	0	1	0	0	0	...
0								
Nepal		1	6	1	2	4	13	...
607								
Pakistan		972	1201	900	668	514	691	...
14314								
Sri Lanka		371	290	197	1086	845	1838	...
4930								
		2006	2007	2008	2009	2010	2011	
2012 \								
Afghanistan		3009	2652	2111	1746	1758	2203	
2635								
Bangladesh		4014	2897	2939	2104	4721	2694	
2640								
Bhutan		10	7	36	865	1464	1879	
1075								
India		33848	28742	28261	29456	34235	27509	
30933								
Iran (Islamic Republic of)		7480	6974	6475	6580	7477	7479	
7534								
Maldives		0	2	1	7	4	3	
1								
Nepal		540	511	581	561	1392	1129	
1185								
Pakistan		13127	10124	8994	7217	6811	7468	
11227								
Sri Lanka		4714	4123	4756	4547	4422	3309	
3338								

	2013	Total
Afghanistan	2004	58639
Bangladesh	3789	65568
Bhutan	487	5876
India	33087	691904
Iran (Islamic Republic of)	11291	175923
Maldives	1	30
Nepal	1308	10222
Pakistan	12603	241600
Sri Lanka	2394	148358

[9 rows x 38 columns]

Exercise: Fetch the data where AreaName is 'Africa' and RegName is 'Southern Africa'. Display the dataframe and find out how many instances are there?

```
df_can[(df_can['Continent']=='Africa') & (df_can['Region']=='Southern Africa')]
```

		Continent	Region	DevName	1980
1981	1982	Africa	Southern Africa	Developing regions	1026
South Africa	\				
1118	781				
Botswana					
1	3				
Namibia					
5	5				
Lesotho		Africa	Southern Africa	Developing regions	1
1	1				
Swaziland					
1	1	Africa	Southern Africa	Developing regions	4
1983	1984				
1985	1986				
...	...				
2005	2006				
2007	2008				
2009	2010				
South Africa	\				
1188	1238				
Botswana					
15	42				
Namibia					
14	16				
Lesotho					
8	7				
Swaziland					
10	3				
2011	2012	2013	Total		
South Africa	959	1243	1240	40568	
Botswana	53	64	76	396	
Namibia	23	24	83	320	

Lesotho	1	0	6	107
Swaziland	13	17	39	188

[5 rows x 38 columns]

Sorting Values of a Dataframe or Series

You can use the `sort_values()` function is used to sort a DataFrame or a Series based on one or more columns. You to specify the column(s) by which you want to sort and the order (ascending or descending). Below is the syntax to use it:- `df.sort_values(col_name, axis=0, ascending=True, inplace=False, ignore_index=False)` `col_name` - the column(s) to sort by. `axis` - axis along which to sort. 0 for sorting by rows (default) and 1 for sorting by columns. `ascending` - to sort in ascending order (True, default) or descending order (False). `inplace` - to perform the sorting operation in-place (True) or return a sorted copy (False, default). `ignore_index` - to reset the index after sorting (True) or keep the original index values (False, default).

Let's sort out dataframe `df_can` on 'Total' column, in descending order to find out the top 5 countries that contributed the most to immigration to Canada.

```
df_can.sort_values(by='Total', ascending=False, axis=0, inplace=True)
top_5 = df_can.head(5)
top_5
```

	Continent \
India	Asia
China	Asia
United Kingdom of Great Britain and Northern Ir...	Europe
Philippines	Asia
Pakistan	Asia

	Region
\	
India	Southern Asia
China	Eastern Asia
United Kingdom of Great Britain and Northern Ir...	Northern Europe
Philippines	South-Eastern Asia
Pakistan	Southern Asia

	DevName
1980 \	
India	Developing regions
8880	
China	Developing regions
5123	

United Kingdom of Great Britain and Northern Ir...	Developed regions			
22045				
Philippines	Developing regions			
6051				
Pakistan	Developing regions			
978				
		1981	1982	
1983 \				
India		8670	8147	
7338				
China		6682	3308	
1863				
United Kingdom of Great Britain and Northern Ir...		24796	20620	
10015				
Philippines		5921	5249	
4562				
Pakistan		972	1201	
900				
		1984	1985	1986
... \				
India		5704	4211	7150
...				
China		1527	1816	1960
...				
United Kingdom of Great Britain and Northern Ir...		10170	9564	9470
...				
Philippines		3801	3150	4166
...				
Pakistan		668	514	691
...				
		2005	2006	
2007 \				
India		36210	33848	
28742				
China		42584	33518	
27642				
United Kingdom of Great Britain and Northern Ir...		7258	7140	
8216				
Philippines		18139	18400	
19837				
Pakistan		14314	13127	
10124				
		2008	2009	
2010 \				
India		28261	29456	
34235				

China	30037	29622
30391		
United Kingdom of Great Britain and Northern Ir...	8979	8876
8724		
Philippines	24887	28573
38617		
Pakistan	8994	7217
6811		
	2011	2012
2013 \		
India	27509	30933
33087		
China	28502	33024
34129		
United Kingdom of Great Britain and Northern Ir...	6204	6195
5827		
Philippines	36765	34315
29544		
Pakistan	7468	11227
12603		
	Total	
India	691904	
China	659962	
United Kingdom of Great Britain and Northern Ir...	551500	
Philippines	511391	
Pakistan	241600	

[5 rows x 38 columns]

Exercise: Find out top 3 countries that contributes the most to immigration to Canda in the year 2010. Display the country names with the immigrant count in this year

```
df_can.sort_values(by='2010', ascending=False, axis=0, inplace=True)
top3_2010 = df_can['2010'].head(3)
top3_2010
```

Philippines	38617
India	34235
China	30391

Name: 2010, dtype: int64

Congratulations! you have learned how to wrangle data with Pandas. You will be using alot of these commands to preprocess the data before its can be used for data visualization.

Thank you for completing this lab!

Author

Alex Aklson

Other Contributors

[Jay Rajasekharan](#), [Ehsan M. Kermani](#), [Slobodan Markovic](#), [Weiqing Wang](#), [Dr. Pooja](#)

Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2023-06-08	2.5	Dr. Pooja	Separated from original lab
2021-05-29	2.4	Weiqing Wang	Fixed typos and code smells.
2021-01-20	2.3	Lakshmi Holla	Changed TOC cell markdown
2020-11-20	2.2	Lakshmi Holla	Changed IBM box URL
2020-11-03	2.1	Lakshmi Holla	Changed URL and info method
2020-08-27	2.0	Lavanya	Moved Lab to course repo in GitLab

© IBM Corporation 2020. All rights reserved.