

# Waffle Charts, Word Clouds, and Regression Plots

Estimated time needed: **40** minutes

## Objectives

After completing this lab you will be able to:

- Create Word cloud and Waffle charts
- Create regression plots with Seaborn library

## Table of Contents

## Import Libraries

```
#Import and setup matplotlib:
%matplotlib inline

import matplotlib as mpl
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches # needed for waffle Charts

mpl.style.use('ggplot') # optional: for ggplot-like style

#Import Primary Modules:
import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library
from PIL import Image # converting images into arrays

#install seaborn and wordcloud
!pip install seaborn wordcloud

#import seaborn
import seaborn as sns

#import wordcloud
import wordcloud

# check for latest version of Matplotlib and seaborn
print ('Matplotlib version: ', mpl.__version__) # >= 2.0.0
print('Seaborn version: ', sns.__version__)
print('WordCloud version: ', wordcloud.__version__)
```

Requirement already satisfied: seaborn in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (0.9.0)  
Requirement already satisfied: wordcloud in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (1.9.3)  
Requirement already satisfied: matplotlib>=1.4.3 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
seaborn) (3.5.3)  
Requirement already satisfied: numpy>=1.9.3 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
seaborn) (1.21.6)  
Requirement already satisfied: pandas>=0.15.2 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
seaborn) (1.3.5)  
Requirement already satisfied: scipy>=0.14.0 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
seaborn) (1.7.3)  
Requirement already satisfied: pillow in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
wordcloud) (9.5.0)  
Requirement already satisfied: cycler>=0.10 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
matplotlib>=1.4.3->seaborn) (0.11.0)  
Requirement already satisfied: fonttools>=4.22.0 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
matplotlib>=1.4.3->seaborn) (4.38.0)  
Requirement already satisfied: kiwisolver>=1.0.1 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
matplotlib>=1.4.3->seaborn) (1.4.4)  
Requirement already satisfied: packaging>=20.0 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
matplotlib>=1.4.3->seaborn) (23.1)  
Requirement already satisfied: pyparsing>=2.2.1 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
matplotlib>=1.4.3->seaborn) (3.0.9)  
Requirement already satisfied: python-dateutil>=2.7 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
matplotlib>=1.4.3->seaborn) (2.8.2)  
Requirement already satisfied: pytz>=2017.3 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
pandas>=0.15.2->seaborn) (2023.3)  
Requirement already satisfied: typing-extensions in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
kiwisolver>=1.0.1->matplotlib>=1.4.3->seaborn) (4.5.0)  
Requirement already satisfied: six>=1.5 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
python-dateutil>=2.7->matplotlib>=1.4.3->seaborn) (1.16.0)  
Matplotlib version: 3.5.3  
Seaborn version: 0.9.0  
WordCloud version: 1.9.3

# Fetching Data

Toolkits: The course heavily relies on *pandas* and *Numpy* for data wrangling, analysis, and visualization. The primary plotting library we will explore in the course is *Matplotlib*.

Dataset: Immigration to Canada from 1980 to 2013 - [International migration flows to and from selected countries - The 2015 revision](#) from United Nation's website

The dataset contains annual data on the flows of international migrants as recorded by the countries of destination. The data presents both inflows and outflows according to the place of birth, citizenship or place of previous / next residence both for foreigners and nationals. In this lab, we will focus on the Canadian Immigration data and use the *already cleaned dataset*. You can refer to the lab on data pre-processing wherein this dataset is cleaned for a quick refresh your Panads skill [Data pre-processing with Pandas](#)

Download the Canadian Immigration dataset and read it into a *pandas* dataframe.

```
df_can = pd.read_csv('https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DV0101EN-SkillsNetwork/Data%20Files/Canada.csv')
```

```
print('Data read into a pandas dataframe!')
```

Data read into a pandas dataframe!

Let's take a look at the first five items in our dataset

```
df_can.head()
```

Country					Continent					Region					DevName	1980
1981 \	039	Afghanistan		Asia		Southern Asia		Developing	regions						16	
10	267	Albania		Europe		Southern Europe		Developed	regions						1	
2	67	Algeria		Africa		Northern Africa		Developing	regions						80	
3	1	American Samoa		Oceania		Polynesia		Developing	regions						0	
4	0	Andorra		Europe		Southern Europe		Developed	regions						0	
1982	1983	1984	1985	...	2005	2006	2007	2008	2009	2010						
2011 \	039	47	71	340	...	3436	3009	2652	2111	1746	1758					
2203	1	0	0	0	...	1223	856	702	560	716	561					
539	2	71	69	63	44	...	3626	4807	3623	4005	5393	4752				
4325																

	2012	2013	Total
0	2635	2004	58639
1	620	603	15699
2	3774	4331	69439
3	0	0	6
4	1	1	15

Let's find out how many entries there are in our dataset

1981 \ Country	Continent	Region	DevName	1980
Afghanistan 39	Asia	Southern Asia	Developing regions	16
Albania 0	Europe	Southern Europe	Developed regions	1
Algeria 67	Africa	Northern Africa	Developing regions	80
American Samoa 1	Oceania	Polynesia	Developing regions	0
Andorra 0	Europe	Southern Europe	Developed regions	0

	1982	1983	1984	1985	1986	...	2005	2006	2007
2008 \ Country						...			
Afghanistan	39	47	71	340	496	...	3436	3009	2652
Albania	0	0	0	0	1	...	1223	856	702
Algeria	71	69	63	44	69	...	3626	4807	3623



Sweden	128	158	187	...	205	139	193	165	167	159
134										

	2012	2013	Total
Country			
Denmark	94	81	3901
Norway	53	59	2327
Sweden	140	140	5866

[3 rows x 38 columns]

Unfortunately, unlike R, `waffle` charts are not built into any of the Python visualization libraries. Therefore, we will learn how to create them from scratch.

**Step 1.** The first step into creating a waffle chart is determining the proportion of each category with respect to the total.

```
# compute the proportion of each category with respect to the total
total_values = df_dsn['Total'].sum()
category_proportions = df_dsn['Total'] / total_values

# print out proportions
pd.DataFrame({"Category Proportion": category_proportions})
```

	Category Proportion
Country	
Denmark	0.322557
Norway	0.192409
Sweden	0.485034

**Step 2.** The second step is defining the overall size of the `waffle` chart.

```
width = 40 # width of chart
height = 10 # height of chart

total_num_tiles = width * height # total number of tiles

print(f'Total number of tiles is {total_num_tiles}.')

Total number of tiles is 400.
```

**Step 3.** The third step is using the proportion of each category to determine its respective number of tiles

```
# compute the number of tiles for each category
tiles_per_category = (category_proportions *
total_num_tiles).round().astype(int)
```

```
# print out number of tiles per category
pd.DataFrame({"Number of tiles": tiles_per_category})
```

	Number of tiles
Country	
Denmark	129
Norway	77
Sweden	194

Based on the calculated proportions, Denmark will occupy 129 tiles of the waffle chart, Norway will occupy 77 tiles, and Sweden will occupy 194 tiles.

**Step 4.** The fourth step is creating a matrix that resembles the waffle chart and populating it.

```
# initialize the waffle chart as an empty matrix
waffle_chart = np.zeros((height, width), dtype = np.uint)

# define indices to loop through waffle chart
category_index = 0
tile_index = 0

# populate the waffle chart
for col in range(width):
    for row in range(height):
        tile_index += 1

        # if the number of tiles populated for the current category is
        # equal to its corresponding allocated tiles...
        if tile_index > sum(tiles_per_category[0:category_index]):
            # ...proceed to the next category
            category_index += 1

        # set the class value to an integer, which increases with
        class
        waffle_chart[row, col] = category_index

print ('Waffle chart populated!')

Waffle chart populated!
```

Let's take a peek at how the matrix looks like.

```
waffle_chart  
array([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2,  
3,  
      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3],  
      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2,  
3,  
      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3]),
```

```

3,      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
3,      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3],
3,      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
3,      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3],
3,      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
3,      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3],
3,      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
3,      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3],
3,      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
3,      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3],
3,      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
3,      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3],
3,      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
3,      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3],
3,      [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
3,      3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3]],
dtype=uint64)

```

As expected, the matrix consists of three categories and the total number of each category's instances matches the total number of tiles allocated to each category.

**Step 5.** Map the waffle chart matrix into a visual.

```


# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()
plt.show()

<Figure size 640x480 with 0 Axes>

```





**Step 6.** Prettify the chart.

```
# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()


# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])
plt.show()

<Figure size 640x480 with 0 Axes>
```



**Step 7.** Create a legend and add it to chart.

```

# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])


# compute cumulative sum of individual categories to match color
schemes between chart and legend
values_cumsum = np.cumsum(df_dsn['Total'])
total_values = values_cumsum[len(values_cumsum) - 1]

# create legend
legend_handles = []
for i, category in enumerate(df_dsn.index.values):
    label_str = category + ' (' + str(df_dsn['Total'][i]) + ')'
    color_val = colormap(float(values_cumsum[i])/total_values)
    legend_handles.append(mpatches.Patch(color=color_val,
label=label_str))

# add legend to chart
plt.legend(handles=legend_handles,
          loc='lower center',
          ncol=len(df_dsn.index.values),
          bbox_to_anchor=(0., -0.2, 0.95, .1)
          )
plt.show()

```

<Figure size 640x480 with 0 Axes>



And there you go! What a good looking *delicious* waffle chart, don't you think?

Now it would be very inefficient to repeat these seven steps every time we wish to create a waffle chart. So let's combine all seven steps into one function called `create_waffle_chart`. This function would take the following parameters as input:

1. **categories:** Unique categories or classes in dataframe.
2. **values:** Values corresponding to categories or classes.
3. **height:** Defined height of waffle chart.
4. **width:** Defined width of waffle chart.
5. **colormap:** Colormap class
6. **value\_sign:** In order to make our function more generalizable, we will add this parameter to address signs that could be associated with a value such as %, \$, and so on. **value\_sign** has a default value of empty string.

```
def create_waffle_chart(categories, values, height, width, colormap,
value_sign=''):

    # compute the proportion of each category with respect to the
    total
    total_values = sum(values)
    category_proportions = [(float(value) / total_values) for value in
values]

    # compute the total number of tiles
    total_num_tiles = width * height # total number of tiles
    print('Total number of tiles is', total_num_tiles)

    # compute the number of tiles for each category
    tiles_per_category = [round(proportion * total_num_tiles) for
proportion in category_proportions]

    # print out number of tiles per category
    for i, tiles in enumerate(tiles_per_category):
        print(df_dsn.index.values[i] + ': ' + str(tiles))

    # initialize the waffle chart as an empty matrix
    waffle_chart = np.zeros((height, width))
```

```

# define indices to loop through waffle chart
category_index = 0
tile_index = 0

# populate the waffle chart
for col in range(width):
    for row in range(height):
        tile_index += 1

        # if the number of tiles populated for the current
category    # is equal to its corresponding allocated tiles...
        if tile_index > sum(tiles_per_category[0:category_index]):
            # ...proceed to the next category
            category_index += 1

        # set the class value to an integer, which increases with
class        waffle_chart[row, col] = category_index

# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

# add dridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])

# compute cumulative sum of individual categories to match color
schemes between chart and legend
values_cumsum = np.cumsum(values)
total_values = values_cumsum[len(values_cumsum) - 1]

# create legend
legend_handles = []
for i, category in enumerate(categories):

```

```

        if value_sign == '%':
            label_str = category + ' (' + str(values[i]) + value_sign
+ ')'
        else:
            label_str = category + ' (' + value_sign + str(values[i])
+ ')'

        color_val = colormap(float(values_cumsum[i])/total_values)
        legend_handles.append(mpatches.Patch(color=color_val,
label=label_str))

    # add legend to chart
    plt.legend(
        handles=legend_handles,
        loc='lower center',
        ncol=len(categories),
        bbox_to_anchor=(0., -0.2, 0.95, .1)
    )
    plt.show()

```

Now to create a waffle chart, all we have to do is call the function `create_waffle_chart`. Let's define the input parameters:

```

width = 40 # width of chart
height = 10 # height of chart

categories = df_dsn.index.values # categories
values = df_dsn['Total'] # corresponding values of categories

colormap = plt.cm.coolwarm # color map class

```

And now let's call our function to create a waffle chart.


```

create_waffle_chart(categories, values, height, width, colormap)

Total number of tiles is 400
Denmark: 129
Norway: 77
Sweden: 194

<Figure size 640x480 with 0 Axes>

```



There seems to be a new Python package for generating waffle charts called [PyWaffle](#), Let's create the same waffle chart with **pywaffle** now

```
#install pywaffle
```

```
!pip install pywaffle
```

```
Requirement already satisfied: pywaffle in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (1.1.1)
Requirement already satisfied: fontawesomefree in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
pywaffle) (6.5.1)
Requirement already satisfied: matplotlib in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
pywaffle) (3.5.3)
Requirement already satisfied: cyclar>=0.10 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
matplotlib->pywaffle) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
matplotlib->pywaffle) (4.38.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
matplotlib->pywaffle) (1.4.4)
Requirement already satisfied: numpy>=1.17 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
matplotlib->pywaffle) (1.21.6)
Requirement already satisfied: packaging>=20.0 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
matplotlib->pywaffle) (23.1)
Requirement already satisfied: pillow>=6.2.0 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
matplotlib->pywaffle) (9.5.0)
Requirement already satisfied: pyparsing>=2.2.1 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
matplotlib->pywaffle) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from
matplotlib->pywaffle) (2.8.2)
```

Requirement already satisfied: typing-extensions in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
kiwisolver>=1.0.1->matplotlib->pywaffle) (4.5.0)


Requirement already satisfied: six>=1.5 in  
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (from  
python-dateutil>=2.7->matplotlib->pywaffle) (1.16.0)

```
#import Waffle from pywaffle
```

```
from pywaffle import Waffle
```

```
#Set up the Waffle chart figure
```

```
fig = plt.figure(FigureClass = Waffle,  
                 rows = 20, columns = 30, #pass the number of rows and  
columns for the waffle  
                 values = df_dsn['Total'], #pass the data to be used  
for display  
                 cmap_name = 'tab20', #color scheme  
                 legend = {'labels': [f"{k} ({v})" for k, v in  
zip(df_dsn.index.values,df_dsn.Total)],  
                           'loc': 'lower left', 'bbox_to_anchor':(0,-  
0.1),'ncol': 3}  
                 #notice the use of list comprehension for creating  
labels  
                 #from index and total of the dataset  
                 )  
  
#Display the waffle chart  
plt.show()
```



**Question:** Create a Waffle chart to display the proportion of China and India total immigrant contribution.

```
# let's create a new dataframe for these three countries
```

```
df_CI = df_can.loc[['China', 'India'], :]
```

```
# let's take a look at our dataframe
```

```
df_CI
```

	Continent	Region	DevName	1980	1981	1982
1983 \ Country						
China	Asia	Eastern Asia	Developing regions	5123	6682	3308
1863						
India	Asia	Southern Asia	Developing regions	8880	8670	8147
7338						
	1984	1985	1986	...	2005	2006
2010 \ Country				...	2007	2008
					2009	
China	1527	1816	1960	...	42584	33518
30391					27642	30037
						29622



India	5704	4211	7150	...	36210	33848	28742	28261	29456
34235									

	2011	2012	2013	Total
Country				
China	28502	33024	34129	659962
India	27509	30933	33087	691904

[2 rows x 38 columns]

*# compute the proportion of each category with respect to the total*

```
total_values = df_CI['Total'].sum()
category_proportions = df_CI['Total'] / total_values
```

*# print out proportions*

```
pd.DataFrame({"Category Proportion": category_proportions})
```

	Category Proportion
Country	
China	0.488186
India	0.511814

```
width = 40 # width of chart
height = 10 # height of chart
```

```
total_num_tiles = width * height # total number of tiles
```

```
print(f'Total number of tiles is {total_num_tiles}.')
```

Total number of tiles is 400.

*# compute the number of tiles for each category*

```
tiles_per_category = (category_proportions *
total_num_tiles).round().astype(int)
```

*# print out number of tiles per category*

```
pd.DataFrame({"Number of tiles": tiles_per_category})
```

	Number of tiles
Country	
China	195
India	205

*# initialize the waffle chart as an empty matrix*

```
waffle_chart = np.zeros((height, width), dtype = np.uint)
```

*# define indices to loop through waffle chart*

```
category_index = 0
tile_index = 0
```

*# populate the waffle chart*

```
for col in range(width):
```

[illegible]


```

# instantiate a new figure object
fig = plt.figure()

# use imshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.imshow(waffle_chart, cmap=colormap)
plt.colorbar()
plt.show()

```

<Figure size 640x480 with 0 Axes>



```

# instantiate a new figure object
fig = plt.figure()

# use imshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.imshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])

# compute cumulative sum of individual categories to match color
schemes between chart and legend
values_cumsum = np.cumsum(df_CI['Total'])
total_values = values_cumsum[len(values_cumsum) - 1]

# create legend


```

```

legend_handles = []
for i, category in enumerate(df_CI.index.values):
    label_str = category + ' (' + str(df_CI['Total'][i]) + ')'
    color_val = colormap(float(values_cumsum[i])/total_values)
    legend_handles.append(mpatches.Patch(color=color_val,
label=label_str))

# add legend to chart
plt.legend(handles=legend_handles,
          loc='lower center',
          ncol=len(df_dsn.index.values),
          bbox_to_anchor=(0., -0.2, 0.95, .1)
        )
plt.show()
<Figure size 640x480 with 0 Axes>

```



## Word Clouds

Word clouds (also known as text clouds or tag clouds) work in a simple way: the more a specific word appears in a source of textual data (such as a speech, blog post, or database), the bigger and bolder it appears in the word cloud.

Luckily, a Python package already exists in Python for generating word clouds. The package, called `wordcloud` was developed by **Andreas Mueller**. You can learn more about the package by following this [link](#).

Let's use this package to learn how to generate a word cloud for a given text document.

First, let's install the package.

```

#import package and its set of stopwords
from wordcloud import WordCloud, STOPWORDS

print ('Wordcloud imported!')

```

```
Wordcloud imported!
```

Word clouds are commonly used to perform high-level analysis and visualization of text data. Accordingly, let's digress from the immigration dataset and work with an example that involves analyzing text data. Let's try to analyze a short novel written by **Lewis Carroll** titled *Alice's Adventures in Wonderland*. Let's go ahead and download a *.txt* file of the novel.

```
import urllib

# # open the file and read it into a variable alice_novel
alice_novel = urllib.request.urlopen('https://cf-courses-
data.s3.us.cloud-object-storage.appdomain.cloud/
IBMDeveloperSkillsNetwork-DV0101EN-SkillsNetwork/Data%20Files/
alice_novel.txt').read().decode("utf-8")
print('Done')
```

Done

Next, let's use the stopwords that we imported from `word_cloud`. We use the function `set` to remove any redundant stopwords.

```
stopwords = set(STOPWORDS)
```

Create a word cloud object and generate a word cloud. For simplicity, let's generate a word cloud using only the first 2000 words in the novel.

```
#if you get attribute error while generating wordcloud, upgrade
Pillow and numpy using below code
%pip install --upgrade Pillow
%pip install --upgrade numpy

#After upgrading, Restart the kernel, delete the above upgrading code
cell and rerun the codes from the start.

Requirement already satisfied: Pillow in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages (9.5.0)
Note: you may need to restart the kernel to use updated packages.
Requirement already satisfied: numpy in
/home/jupyterlab/conda/envs/python/lib/python3.7/site-packages
(1.21.6)
Note: you may need to restart the kernel to use updated packages.


# instantiate a word cloud object
alice_wc = WordCloud()

# generate the word cloud
alice_wc.generate(alice_novel)

<wordcloud.wordcloud.WordCloud at 0x7f8c52079f50>
```

Awesome! Now that the word cloud is created, let's visualize it.

```
# display the word cloud
plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```




Interesting! So in the first 2000 words in the novel, the most common words are **Alice**, **said**, **little**, **Queen**, and so on. Let's resize the cloud so that we can see the less frequent words a little better.

```
fig = plt.figure(figsize=(14, 18))

# display the cloud
plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```







Excellent! This looks really interesting! Another cool thing you can implement with the `word_cloud` package is superimposing the words onto a mask of any shape. Let's use a mask of Alice and her rabbit. We already created the mask for you, so let's go ahead and download it and call it `alice_mask.png`.

```
#save mask to alice_mask
alice_mask = np.array(Image.open(urllib.request.urlopen('https://cf-
courses-data.s3.us.cloud-object-storage.appdomain.cloud/
IBMDeveloperSkillsNetwork-DV0101EN-SkillsNetwork/labs/Module%204/
images/alice_mask.png')))
```

Let's take a look at how the mask looks like.

```
fig = plt.figure(figsize=(14, 18))

plt.imshow(alice_mask, cmap=plt.cm.gray, interpolation='bilinear')
plt.axis('off')
plt.show()
```





Shaping the word cloud according to the mask is straightforward using `word_cloud` package. For simplicity, we will continue using the first 2000 words in the novel.

```
# instantiate a word cloud object
alice_wc = WordCloud(background_color='white', max_words=2000,
mask=alice_mask, stopwords=stopwords)

# generate the word cloud
alice_wc.generate(alice_novel)

# display the word cloud
fig = plt.figure(figsize=(14, 18))
```

```
plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```



Really impressive!

Unfortunately, our immigration data does not have any text data, but where there is a will there is a way. Let's generate sample text data from our immigration dataset, say text data of 90 words.

Let's recall how our data looks like.

```
df_can.head()
```

	Continent			Region			DevName	1980	
1981 \ Country									
Afghanistan 39	Asia		Southern Asia		Developing regions			16	
Albania 0	Europe		Southern Europe		Developed regions			1	
Algeria 67	Africa		Northern Africa		Developing regions			80	
American Samoa 1	Oceania		Polynesia		Developing regions			0	
Andorra 0	Europe		Southern Europe		Developed regions			0	
	1982	1983	1984	1985	1986	...	2005	2006	2007
2008 \ Country						...			
Afghanistan 2111	39	47	71	340	496	...	3436	3009	2652
Albania 560	0	0	0	0	1	...	1223	856	702
Algeria 4005	71	69	63	44	69	...	3626	4807	3623
American Samoa 0	0	0	0	0	0	...	0	1	0
Andorra 0	0	0	0	0	2	...	0	1	1
	2009	2010	2011	2012	2013	Total			
Country									
Afghanistan	1746	1758	2203	2635	2004	58639			
Albania	716	561	539	620	603	15699			
Algeria	5393	4752	4325	3774	4331	69439			
American Samoa	0	0	0	0	0	6			
Andorra	0	0	0	1	1	15			

```
[5 rows x 38 columns]
```

And what was the total immigration from 1980 to 2013?

```
total_immigration = df_can['Total'].sum()
total_immigration
```

```
6409153
```

Using countries with single-word names, let's duplicate each country's name based on how much they contribute to the total immigration.

```
max_words = 90
word_string = ''
for country in df_can.index.values:
    # check if country's name is a single-word name
    if country.count(" ") == 0:
        repeat_num_times = int(df_can.loc[country, 'Total'] /
                                total_immigration * max_words)
        word_string = word_string + ((country + ' ') *
                                      repeat_num_times)

# display the generated text
word_string

'China China China China China China China China China Colombia Egypt
France Guyana Haiti India India India India India India India India
India Jamaica Lebanon Morocco Pakistan Pakistan Pakistan Philippines
Philippines Philippines Philippines Philippines Philippines
Philippines Poland Portugal Romania '
```

We are not dealing with any stopwords here, so there is no need to pass them when creating the word cloud.

```
# create the word cloud
wordcloud = WordCloud(background_color='white').generate(word_string)

print('Word cloud created!')

Word cloud created!

# display the cloud
plt.figure(figsize=(14, 18))

plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()
```



According to the above word cloud, it looks like the majority of the people who immigrated came from one of 15 countries that are displayed by the word cloud. One cool visual that you could build, is perhaps using the map of Canada and a mask and superimposing the word cloud on top of the map of Canada. That would be an interesting visual to build!

## Plotting with Seaborn

Seaborn is a Python visualization library based on matplotlib. It provides a high-level interface for drawing attractive statistical graphics. You can learn more about *seaborn* by following this [link](#) and more about *seaborn* regression plots by following this [link](#).

In lab *Pie Charts, Box Plots, Scatter Plots, and Bubble Plots*, we learned how to create a scatter plot and then fit a regression line. It took ~20 lines of code to create the scatter plot along with the regression fit. In this final section, we will explore *seaborn* and see how efficient it is to create regression lines and fits using this library!

### Categorical Plots


In our data 'df\_can', let's find out how many continents are mentioned

```
df_can['Continent'].unique()
array(['Asia', 'Europe', 'Africa', 'Oceania',
      'Latin America and the Caribbean', 'Northern America'],
      dtype=object)
```

## countplot

A count plot can be thought of as a histogram across a categorical, instead of quantitative, variable. Let's find the count of Continents in the data 'df\_can' using countplot on 'Continent'


```
sns.countplot(x='Continent', data=df_can)
<AxesSubplot:xlabel='Continent', ylabel='count'>
```



The labels on the x-axis doesnot look as expected.Let's try to replace the 'Latin America and the Caribbean' with and "L-America", 'Northern America' with "N-America", and change the figure size and then display the plot again

```
df_can1 = df_can.replace('Latin America and the Caribbean', 'L-America')
df_can1 = df_can1.replace('Northern America', 'N-America')

plt.figure(figsize=(15, 10))
sns.countplot(x='Continent', data=df_can1)
<AxesSubplot:xlabel='Continent', ylabel='count'>
```




Much better!

## Barplot

**This plot will perform the Groupby on a categorical variable and plot aggregated values, with confidence intervals.** Let's plot the total immigrants Continent-wise

```
plt.figure(figsize=(15, 10))
sns.barplot(x='Continent', y='Total', data=df_can1)
<AxesSubplot:xlabel='Continent', ylabel='Total'>
```



You can verify the values by performing the groupby on the Total and Continent for mean()

```
df_Can2=df_can1.groupby('Continent')['Total'].mean()
df_Can2
```

```
Continent
Africa      11462.000000
Asia        67710.081633
Europe      32812.720930
L-America   23186.303030
N-America   120571.000000
Oceania      3941.000000
Name: Total, dtype: float64
```

Create a new dataframe that stores that total number of landed immigrants to Canada per year from 1980 to 2013.

## Regression Plot

With *seaborn*, generating a regression plot is as simple as calling the **regplot** function.

```
years = list(map(str, range(1980, 2014)))
# we can use the sum() method to get the total population per year
```



```
df_tot = pd.DataFrame(df_can[years].sum(axis=0))

# change the years to type float (useful for regression later on)
df_tot.index = map(float, df_tot.index)

# reset the index to put in back in as a column in the df_tot
dataframe
df_tot.reset_index(inplace=True)


# rename columns
df_tot.columns = ['year', 'total']

# view the final dataframe
df_tot.head()
```

	year	total
0	1980.0	99137
1	1981.0	110563
2	1982.0	104271
3	1983.0	75550
4	1984.0	73417


```
#seaborn is already imported at the start of this lab
sns.regplot(x='year', y='total', data=df_tot)

<AxesSubplot:xlabel='year', ylabel='total'>
```




This is not magic; it is *seaborn*! You can also customize the color of the scatter plot and regression line. Let's change the color to green.

```
sns.regplot(x='year', y='total', data=df_tot, color='green')  
plt.show()
```




You can always customize the marker shape, so instead of circular markers, let's use +.

```
ax = sns.regplot(x='year', y='total', data=df_tot, color='green',  
marker='+')  
plt.show()
```



Let's blow up the plot a little so that it is more appealing to the sight.

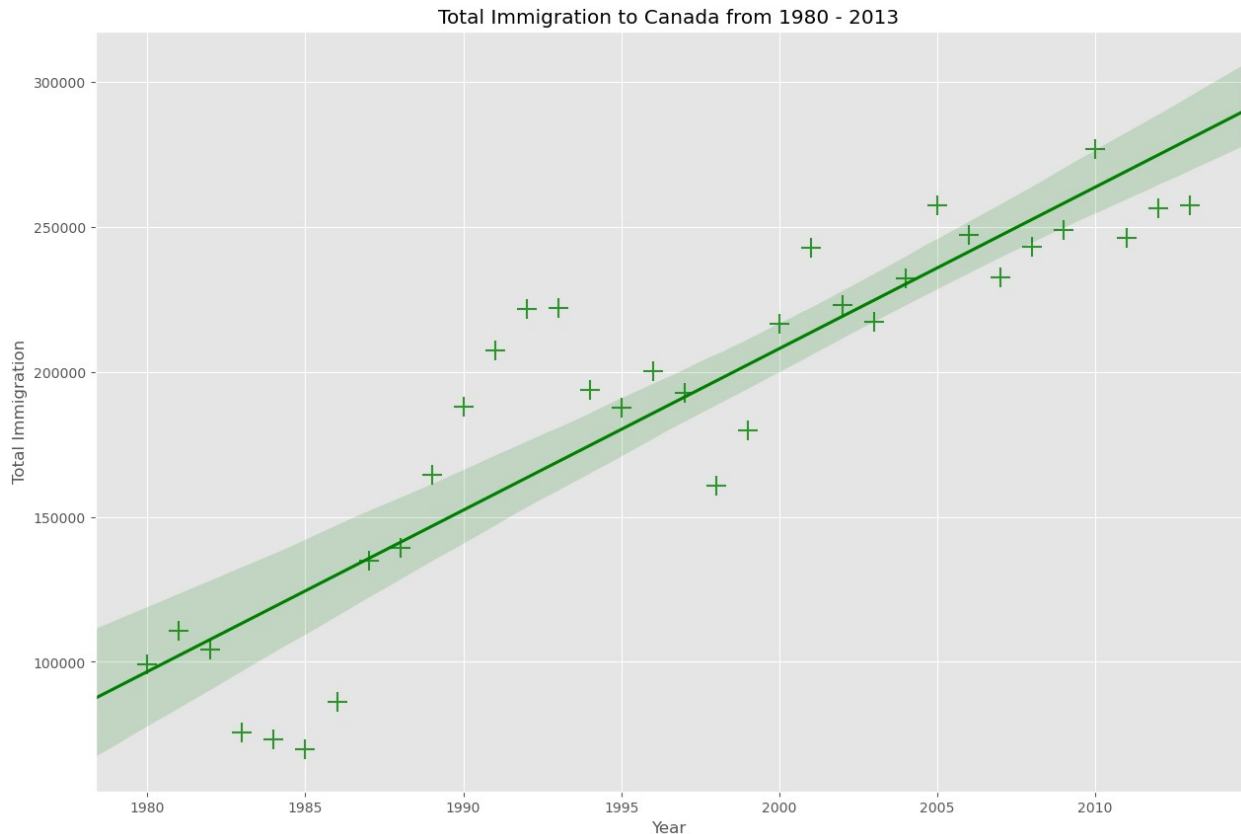
```
plt.figure(figsize=(15, 10))
sns.regplot(x='year', y='total', data=df_tot, color='green',
marker='+')
plt.show()
```



And let's increase the size of markers so they match the new size of the figure, and add a title and x- and y-labels.

```
plt.figure(figsize=(15, 10))
ax = sns.regplot(x='year', y='total', data=df_tot, color='green',
marker='+', scatter_kws={'s': 200})

ax.set(xlabel='Year', ylabel='Total Immigration') # add x- and y-
labels
ax.set_title('Total Immigration to Canada from 1980 - 2013') # add
title
plt.show()
```




And finally increase the font size of the tickmark labels, the title, and the x- and y-labels so they don't feel left out!

```
plt.figure(figsize=(15, 10))

sns.set(font_scale=1.5)

ax = sns.regplot(x='year', y='total', data=df_tot, color='green',
                 marker='+', scatter_kws={'s': 200})
ax.set(xlabel='Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
plt.show()
```




Amazing! A complete scatter plot with a regression fit with 5 lines of code only. Isn't this really amazing?

If you are not a big fan of the purple background, you can easily change the style to a white plain background.

```
plt.figure(figsize=(15, 10))

sns.set(font_scale=1.5)
sns.set_style('ticks') # change background to white background

ax = sns.regplot(x='year', y='total', data=df_tot, color='green',
marker='+', scatter_kws={'s': 200})
ax.set_xlabel('Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
plt.show()
```




Or to a white background with gridlines.

```
plt.figure(figsize=(15, 10))

sns.set(font_scale=1.5)
sns.set_style('whitegrid')

ax = sns.regplot(x='year', y='total', data=df_tot, color='green',
                 marker='+', scatter_kws={'s': 200})
ax.set(xlabel='Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
plt.show()
```





**Question:** Use seaborn to create a scatter plot with a regression line to visualize the total immigration from Denmark, Sweden, and Norway to Canada from 1980 to 2013.

```
# let's create a new dataframe for these three countries
df_countries = df_can.loc[['Denmark', 'Norway', 'Sweden'],
years].transpose()
```

```
# let's take a look at our dataframe
df_countries
```

Country	Denmark	Norway	Sweden
1980	272	116	281
1981	293	77	308
1982	299	106	222
1983	106	51	176
1984	93	31	128
1985	73	54	158
1986	93	56	187
1987	109	80	198
1988	129	73	171
1989	129	76	182
1990	118	83	130
1991	111	103	167
1992	158	74	179
1993	186	92	203

1994	93	60	192
1995	111	65	176
1996	70	70	161
1997	83	104	151
1998	63	31	123
1999	81	36	170
2000	93	56	138
2001	81	78	184
2002	70	74	149
2003	89	77	161
2004	89	73	129
2005	62	57	205
2006	101	53	139
2007	97	73	193
2008	108	66	165
2009	81	75	167
2010	92	46	159
2011	93	49	134
2012	94	53	140
2013	81	59	140

```
# create df_total by summing across three countries for each year
df_total = pd.DataFrame(df_countries.sum(axis=1))
```

```
# reset index in place
df_total.reset_index(inplace=True)
```


```
# rename columns
df_total.columns = ['year', 'total']
```

```
# change column year from string to int to create scatter plot
df_total['year'] = df_total['year'].astype(int)
```

```
### type your answer here
plt.figure(figsize=(15, 10))
```

```
sns.set(font_scale=1.5)
sns.set_style('whitegrid')
```

```
ax = sns.regplot(x='year', y='total', data=df_tot, color='green',
marker='+', scatter_kws={'s': 200})
ax.set(xlabel='Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
plt.show()
```



Thank you for completing this lab!

## Author

Alex Aklson Dr. Pooja

## Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2023-07-07	2.7	Dr. Pooja	wordcloud, sns, pip, pywaffle issue resolved
2023-06-11	2.6	Dr. Pooja	Clean data link, pywaffle, Categorical plots included
2021-05-19	2.3	Weiqing Wang	Fixed typos and code spells
2021-01-21	2.2	Lakshmi Holla	Updated TOC markdown cell
2020-11-03	2.1	Lakshmi Holla	Changed URL of excel file
2020-08-27	2.0	Lavanya	Moved lab to course repo in GitLab

© IBM Corporation 2020. All rights reserved.